

# Deep Learning for NLP

## Lecture 3: Training as Optimization and (Neural) Language Models

**Dr. Mohsen Mesgar**

**Ubiquitous Knowledge Processing Lab (UKP Lab)**

# This lecture

- ▶ training as optimization
- ▶ backpropagation
- ▶ language modeling

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples

# Recall

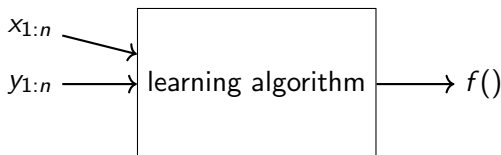
- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels

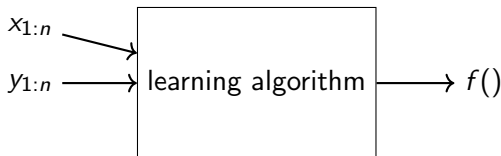
# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels



# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels



- ▶ how to measure if  $f()$  works accurately?



# Loss Function

- ▶ *Loss function*  $L(y, \hat{y})$ : It quantifies the loss suffered when predicting  $\hat{y}$  while the true label is  $y$

# Loss Function

- ▶ *Loss function*  $L(y, \hat{y})$ : It quantifies the loss suffered when predicting  $\hat{y}$  while the true label is  $y$
- ▶ given a labeled training set  $(x_{1:n}, y_{1:n})$ , a per-instance loss function  $L$  and a parameterized function  $f(x; \Theta)$ , we define the corpus-wide loss with respect to the parameters as the average loss over all training examples:

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n L(\hat{y}_i, y_i),$$

$$\hat{y}_i = f(x_i; \Theta)$$

# Loss Function

- ▶ *Loss function*  $L(y, \hat{y})$ : It quantifies the loss suffered when predicting  $\hat{y}$  while the true label is  $y$
- ▶ given a labeled training set  $(x_{1:n}; y_{1:n})$ , a per-instance loss function  $L$  and a parameterized function  $f(x; \Theta)$ , we define the corpus-wide loss with respect to the parameters as the average loss over all training examples:

$$\mathcal{L}(\theta) = \frac{1}{n} \sum_{i=1}^n L(\hat{y}_i, y_i),$$

$$\hat{y}_i = f(x_i; \Theta)$$

- ▶ in this view the training examples are fixed and the values of the parameters determine the loss

# Training as Optimization

- ▶ the goal of the training algorithm is then to set the values of the parameters such that the value of  $\mathcal{L}$  is minimized

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i)$$

# Common Loss Functions

- ▶ Hinge (binary)
  - ▶ for binary classification problems

# Common Loss Functions

- ▶ Hinge (binary)
  - ▶ for binary classification problems
  - ▶ the classifier's output is a single scalar  $\hat{y}$  and the true output  $y$  is in  $\{-1, +1\}$

# Common Loss Functions

- ▶ Hinge (binary)
  - ▶ for binary classification problems
  - ▶ the classifier's output is a single scalar  $\hat{y}$  and the true output  $y$  is in  $\{-1, +1\}$
  - ▶ the inference rule is  $\text{prediction} = \text{sign}(\hat{y})$ , and a classification is considered correct if  $y \cdot \tilde{y} > 0$

# Common Loss Functions

- ▶ Hinge (binary)
  - ▶ for binary classification problems
  - ▶ the classifier's output is a single scalar  $\hat{y}$  and the true output  $y$  is in  $\{-1, +1\}$
  - ▶ the inference rule is  $\text{prediction} = \text{sign}(\hat{y})$ , and a classification is considered correct if  $y \cdot \tilde{y} > 0$
  - ▶ per-instance loss:

$$L_{\text{hinge}(\text{binary})}(\hat{y}, y) = \max(0, 1 - y \cdot \hat{y})$$



# Common Loss Functions

- ▶ Hinge (multi-class)
  - ▶ let  $\hat{y} = \hat{y}_{[1]}, \hat{y}_{[2]}, \dots, \hat{y}_{[n]}$  be the model's output vector, and  $y$  be the one-hot vector for the correct output class

# Common Loss Functions

- ▶ Hinge (multi-class)
  - ▶ let  $\hat{y} = \hat{y}_{[1]}, \hat{y}_{[2]}, \dots, \hat{y}_{[n]}$  be the model's output vector, and  $y$  be the one-hot vector for the correct output class
  - ▶ the inference rule is defined as selecting the class with the highest score prediction =  $\operatorname{argmax}_i \hat{y}_{[i]}$

# Common Loss Functions

- ▶ Hinge (multi-class)
  - ▶ let  $\hat{y} = \hat{y}_{[1]}, \hat{y}_{[2]}, \dots, \hat{y}_{[n]}$  be the model's output vector, and  $y$  be the one-hot vector for the correct output class
  - ▶ the inference rule is defined as selecting the class with the highest score prediction =  $\operatorname{argmax}_i \hat{y}_{[i]}$
  - ▶ if  $t$  is the correct class and  $k$  is the highest scoring class such that  $k \neq t$  then loss is

$$L_{\text{hinge}(\text{multiclass})}(\hat{y}, y) = \max(0, 1 - (\hat{y}_{[t]} - \hat{y}_{[k]}))$$

# Common Loss Functions

- ▶ Log loss
  - ▶ can be seen as a “soft” version of the hinge loss with an infinite margin

# Common Loss Functions

- ▶ Log loss
  - ▶ can be seen as a “soft” version of the hinge loss with an infinite margin
  - ▶ loss

$$L_{\log}(\hat{y}, y) = \log(1 + \exp(-(\hat{y}_{[t]} - \hat{y}_{[k]})))$$

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$



# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ▶ then  $\hat{y} = \sigma(\tilde{y}) = P(y = 1|x)$

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ▶ then  $\hat{y} = \sigma(\tilde{y}) = P(y = 1|x)$
- ▶ inference rule: prediction= 0 if  $\hat{y} < 0.5$  and prediction= 1 if  $\hat{y} \geq 0.5$

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ▶ then  $\hat{y} = \sigma(\tilde{y}) = P(y = 1|x)$
- ▶ inference rule: prediction= 0 if  $\hat{y} < 0.5$  and prediction= 1 if  $\hat{y} \geq 0.5$
- ▶ loss

$$L_{\text{logistic}}(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ▶ then  $\hat{y} = \sigma(\tilde{y}) = P(y = 1|x)$
- ▶ inference rule: prediction= 0 if  $\hat{y} < 0.5$  and prediction= 1 if  $\hat{y} \geq 0.5$
- ▶ loss

$$L_{\text{logistic}}(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

- ▶ is useful for estimating class conditional probability for a binary classification problem

# Common Loss Functions

- ▶ binary cross entropy (logistic loss)
  - ▶ used for binary classification with conditional probability outputs.
  - ▶ assumes a set of two target classes labeled 0 and 1, with a correct label  $y \in \{0, 1\}$
  - ▶ the model's output  $\tilde{y}$  is transformed using the sigmoid function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

- ▶ then  $\hat{y} = \sigma(\tilde{y}) = P(y = 1|x)$
- ▶ inference rule: prediction= 0 if  $\hat{y} < 0.5$  and prediction= 1 if  $\hat{y} \geq 0.5$
- ▶ loss

$$L_{\text{logistic}}(\hat{y}, y) = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

- ▶ is useful for estimating class conditional probability for a binary classification problem
- ▶ we assume that the output layer is transformed using sigmoid

# Training as Optimization

- ▶ the goal of the training algorithm is then to set the values of the parameters such that the value of  $\mathcal{L}$  is minimized

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i)$$

# Training as Optimization

- ▶ the goal of the training algorithm is then to set the values of the parameters such that the value of  $\mathcal{L}$  is minimized

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i)$$

- ▶ the above optimization attempts to minimize the loss at all costs, which may result in overfitting the training data



# Regularization

- ▶ to counter that we define a regularization term  $R(\Theta)$  taking as input the parameters and returning a scalar that reflect their “complexity,” which we want to keep low

# Regularization

- ▶ to counter that we define a regularization term  $R(\Theta)$  taking as input the parameters and returning a scalar that reflect their “complexity,” which we want to keep low
- ▶ training

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i)$$

# Regularization

- ▶ to counter that we define a regularization term  $R(\Theta)$  taking as input the parameters and returning a scalar that reflect their “complexity,” which we want to keep low
- ▶ training

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i)$$

- ▶ training with regularization

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \mathcal{L}(\Theta) = \operatorname{argmin}_{\Theta} \left( \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i) + \lambda R(\Theta) \right)$$

# Regularization

- intuitively we would like to drive the learner toward natural solutions, in which it is OK to mis-classify a few examples if they don't fit well with the rest

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \left( \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i) + \lambda R(\Theta) \right)$$

# Regularization

- ▶ intuitively we would like to drive the learner toward natural solutions, in which it is OK to mis-classify a few examples if they don't fit well with the rest

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \left( \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i) + \lambda R(\Theta) \right)$$

- ▶ regularization term considers the parameter values, and scores their complexity

# Regularization

- ▶ intuitively we would like to drive the learner toward natural solutions, in which it is OK to mis-classify a few examples if they don't fit well with the rest

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} \left( \frac{1}{n} \sum_{i=1}^n L(\hat{y}, y_i) + \lambda R(\Theta) \right)$$

- ▶ regularization term considers the parameter values, and scores their complexity
- ▶ in practice the regularizers equate complexity with large weights and work to keep the parameter values low

# Common Regularization Functions

- ▶  $L_2$  regularization (a.k.a. gaussian prior or weight decay): It keeps the sum of the squares of the parameter values low

$$R_{L_2}(W) = \|W\|_2^2 = \sum_{i,j} (W_{[i,j]})^2$$

# Common Regularization Functions

- ▶  $L_2$  regularization (a.k.a. gaussian prior or weight decay): It keeps the sum of the squares of the parameter values low

$$R_{L_2}(W) = ||W||_2^2 = \sum_{i,j} (W_{[i,j]})^2$$

- ▶ the learner will prefer to decrease the value of one parameter with high weight by 1 than to decrease the value of ten parameters that already have relatively low weights by 0.1 each



# Common Regularization Functions

- ▶  $L_1$  regularization (a.k.a. sparse prior or lasso): It keeps the sum of the absolute values of the parameters low

$$R_{L_1}(W) = \|W\|_1 = \sum_{i,j} |W_{[i,j]}|$$

# Common Regularization Functions

- ▶  $L_1$  regularization (a.k.a. sparse prior or lasso): It keeps the sum of the absolute values of the parameters low

$$R_{L_1}(W) = \|W\|_1 = \sum_{i,j} |W_{[i,j]}|$$

- ▶ the learner will prefer to decrease all the non-zero parameter values toward zero

# Common Regularization Functions

- ▶ Elastic-Net: combines both  $L_1$  and  $L_2$  regularization

$$R_{\text{elastic-net}}(W) = \gamma_1 R_{L_1}(W) + \gamma_2 R_{L_2}(W)$$

# Common Regularization Functions

- ▶ Elastic-Net: combines both  $L_1$  and  $L_2$  regularization

$$R_{\text{elastic-net}}(W) = \gamma_1 R_{L_1}(W) + \gamma_2 R_{L_2}(W)$$

- ▶ dropout: will be discussed later

# Training as Optimization

- ▶ we learned that the goal of training is to minimize a loss function (and a regularization term)

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} (\mathcal{L}(\hat{y}, y) + \lambda R(\Theta))$$

# Training as Optimization

- ▶ we learned that the goal of training is to minimize a loss function (and a regularization term)

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} (\mathcal{L}(\hat{y}, y) + \lambda R(\Theta))$$

- ▶ training = Solving an optimization problem

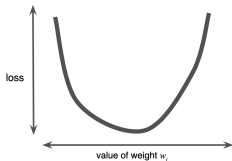
# Training as Optimization

- ▶ we learned that the goal of training is to minimize a loss function (and a regularization term)

$$\hat{\Theta} = \operatorname{argmin}_{\Theta} (\mathcal{L}(\hat{y}, y) + \lambda R(\Theta))$$

- ▶ training = Solving an optimization problem
- ▶ how to find parameter values that minimize loss?

# Gradient-based Optimization

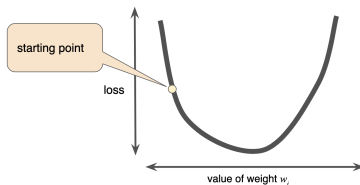


(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)



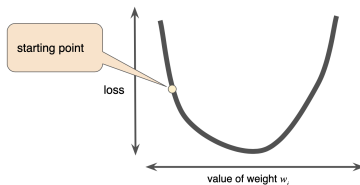
# Gradient-based Optimization



(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)

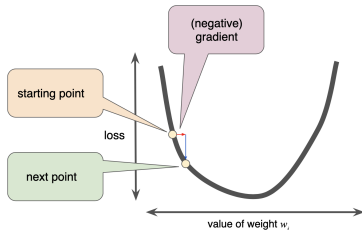
# Gradient-based Optimization



(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)

# Gradient-based Optimization



(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)

- ▶ we repeatedly compute an estimate of the loss over the training set

# Gradient-based Optimization

- ▶ we repeatedly compute an estimate of the loss over the training set
- ▶ we compute the gradients of the parameters with respect to the loss estimate

# Gradient-based Optimization

- ▶ we repeatedly compute an estimate of the loss over the training set
- ▶ we compute the gradients of the parameters with respect to the loss estimate
- ▶ we move the parameter values in the opposite directions of the gradient

# (Online) Stochastic Gradient Descent

---

**Algorithm 2.1** Online stochastic gradient descent training.

---

*Input:*

- Function  $f(\mathbf{x}; \Theta)$  parameterized with parameters  $\Theta$ .
- Training set of inputs  $\mathbf{x}_1, \dots, \mathbf{x}_n$  and desired outputs  $y_1, \dots, y_n$ .
- Loss function  $L$ .

- 
- 1: **while** stopping criteria not met **do**
  - 2:     Sample a training example  $\mathbf{x}_i, y_i$
  - 3:     Compute the loss  $L(f(\mathbf{x}_i; \Theta), y_i)$
  - 4:      $\hat{\mathbf{g}} \leftarrow$  gradients of  $L(f(\mathbf{x}_i; \Theta), y_i)$  w.r.t  $\Theta$
  - 5:      $\Theta \leftarrow \Theta - \eta_t \hat{\mathbf{g}}$
  - 6: **return**  $\Theta$
- 

(Taken from: *Neural Network Methods for Natural Language Processing*, Yoav Goldberg)

# (Minibatch) Stochastic Gradient Descent

---

**Algorithm 2.2** Minibatch stochastic gradient descent training.

---

*Input:*

- Function  $f(\mathbf{x}; \Theta)$  parameterized with parameters  $\Theta$ .
- Training set of inputs  $\mathbf{x}_1, \dots, \mathbf{x}_n$  and desired outputs  $y_1, \dots, y_n$ .
- Loss function  $L$ .

---

```
1: while stopping criteria not met do
2:   Sample a minibatch of  $m$  examples  $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)\}$ 
3:    $\hat{\mathbf{g}} \leftarrow 0$ 
4:   for  $i = 1$  to  $m$  do
5:     Compute the loss  $L(f(\mathbf{x}_i; \Theta), y_i)$ 
6:      $\hat{\mathbf{g}} \leftarrow \hat{\mathbf{g}} + \text{gradients of } \frac{1}{m}L(f(\mathbf{x}_i; \Theta), y_i) \text{ w.r.t } \Theta$ 
7:    $\Theta \leftarrow \Theta - \eta_t \hat{\mathbf{g}}$ 
8: return  $\Theta$ 
```

---

(Taken from: *Neural Network Methods for Natural Language Processing*, Yoav Goldberg)



# Why Minibatch SGD?

- ▶ it's not expensive: while computing loss function gradient on all training set can be computationally expensive

# Why Minibatch SGD?

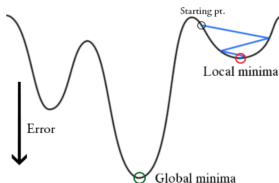
- ▶ it's not expensive: while computing loss function gradient on all training set can be computationally expensive
- ▶ it converges faster to a good solution than full-batch learning, in which we use all training set to compute gradient

# Why Minibatch SGD?

- ▶ it's not expensive: while computing loss function gradient on all training set can be computationally expensive
- ▶ it converges faster to a good solution than full-batch learning, in which we use all training set to compute gradient
- ▶ smaller mini-batch sizes lead often to better solutions (generalize better)

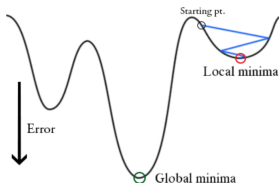
# Stochastic Gradient Descent (SGD)

- ▶ gradient descent does not always lead to best solutions



# Stochastic Gradient Descent (SGD)

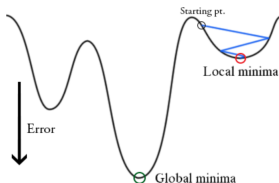
- ▶ gradient descent does not always lead to best solutions



- ▶ why?

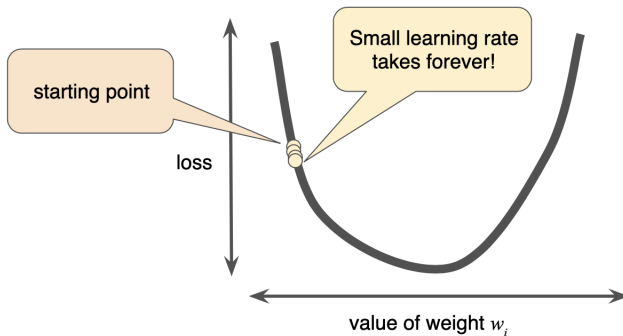
# Stochastic Gradient Descent (SGD)

- ▶ gradient descent does not always lead to best solutions



- ▶ why?
- ▶ SGD is sensitive to the learning rate and initial parameter values (starting point)

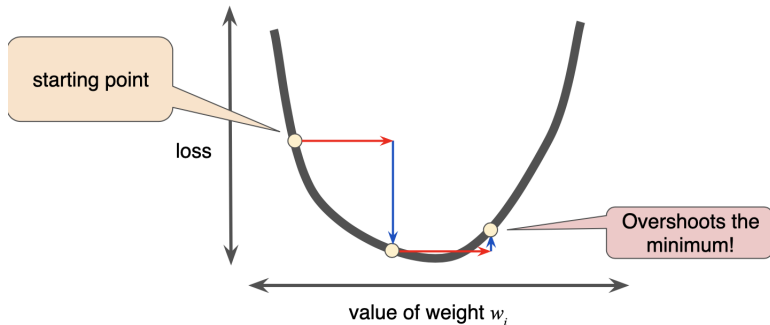
# Small Learning Rate



(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)

# Large Learning Rate

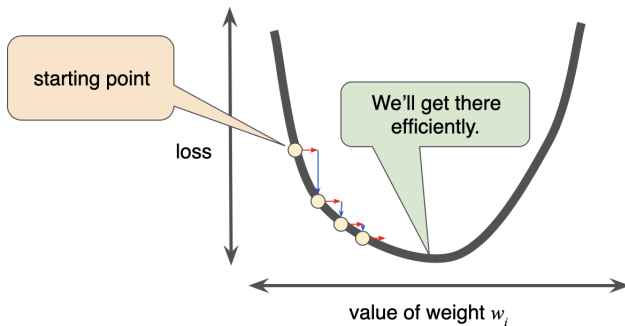


(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)



# Adaptive Learning Rate



(Taken from:

<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>)

# Stochastic Gradient Descent (SGD)

- ▶ Use adaptive learning rate algorithms
  - ▶ AdaGrad [Duchi et al., 2011],
  - ▶ AdaDelta [Zeiler, 2012],
  - ▶ RMSProp [Tieleman and Hinton, 2012],
  - ▶ Adam [Kingma and Ba, 2014]

# Stochastic Gradient Descent (SGD)

- ▶ Use adaptive learning rate algorithms
  - ▶ AdaGrad [Duchi et al., 2011],
  - ▶ AdaDelta [Zeiler, 2012],
  - ▶ RMSProp [Tieleman and Hinton, 2012],
  - ▶ Adam [Kingma and Ba, 2014]
- ▶ these methods train usually faster than SGD

# Stochastic Gradient Descent (SGD)

- ▶ Use adaptive learning rate algorithms
  - ▶ AdaGrad [Duchi et al., 2011],
  - ▶ AdaDelta [Zeiler, 2012],
  - ▶ RMSProp [Tieleman and Hinton, 2012],
  - ▶ Adam [Kingma and Ba, 2014]
- ▶ these methods train usually faster than SGD
- ▶ found solution is often not as good as that by SGD → First train with Adam, fine-tune with SGD

# Stochastic Gradient Descent (SGD)

- ▶ Use adaptive learning rate algorithms
  - ▶ AdaGrad [Duchi et al., 2011],
  - ▶ AdaDelta [Zeiler, 2012],
  - ▶ RMSProp [Tieleman and Hinton, 2012],
  - ▶ Adam [Kingma and Ba, 2014]
- ▶ these methods train usually faster than SGD
- ▶ found solution is often not as good as that by SGD → First train with Adam, fine-tune with SGD
- ▶ they use different initial parameter values in different runs of experiments and report the average of scores

# Backpropagation

- ▶ a fancy name for a recursive algorithm that computes the derivatives of a nested functions using the chain rule, while caching intermediary derivatives

# Backpropagation

- ▶ a fancy name for a recursive algorithm that computes the derivatives of a nested functions using the chain rule, while caching intermediary derivatives
- ▶ chain rule: Assume  $y = f(g(x))$

$$\frac{\partial y}{\partial x} = \frac{\partial f}{\partial g} \times \frac{\partial g}{\partial x} = \frac{\partial f}{\partial g} \frac{\partial g}{\partial x}$$

# Backpropagation

- ▶ consists of two steps



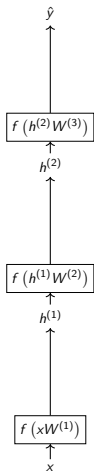
# Backpropagation

- ▶ consists of two steps
  - ▶ forward pass → use current parameter values to compute the loss value

# Backpropagation

- ▶ consists of two steps
  - ▶ forward pass → use current parameter values to compute the loss value
  - ▶ backward pass → use the gradient of the loss to update the parameter values

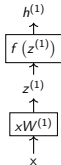
# Model: Multilayer Perceptron (MLP)



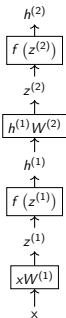
# Backpropagation: Forward Pass



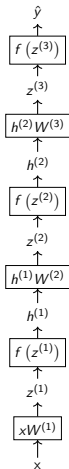
# Backpropagation: Forward Pass



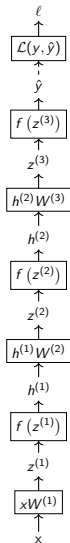
# Backpropagation: Forward Pass



# Backpropagation: Forward Pass

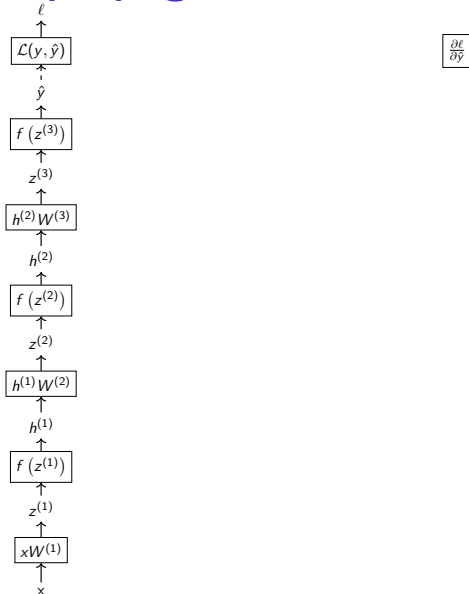


# Backpropagation: Forward Pass

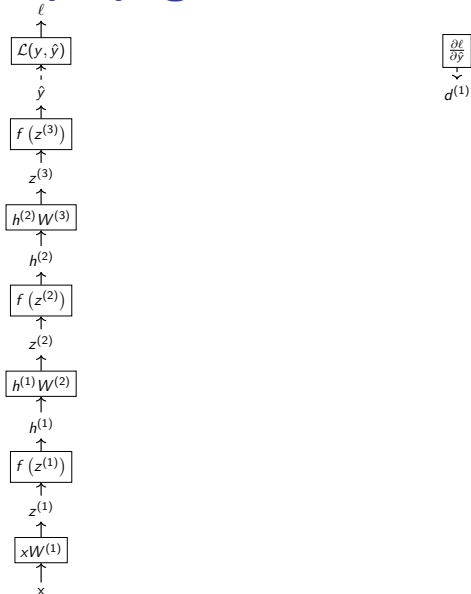




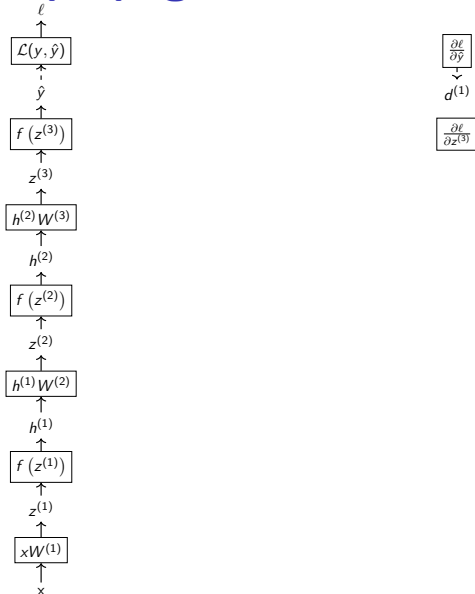
# Backpropagation: Backward Pass



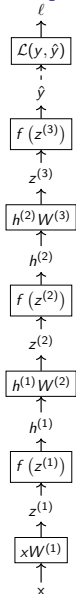
# Backpropagation: Backward Pass



# Backpropagation: Backward Pass



# Backpropagation: Backward Pass



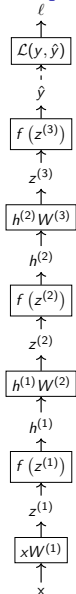
$$\frac{\partial \ell}{\partial \hat{y}}$$

$$\downarrow$$

$$d^{(1)}$$

$$\frac{\partial \ell}{\partial z^{(3)}} = \frac{\partial \ell}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}}$$

# Backpropagation: Backward Pass



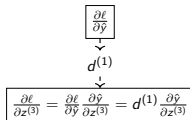
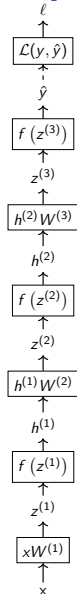
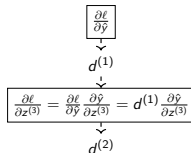


Diagram illustrating the backward pass of a neural network. The error gradient  $\frac{\partial \mathcal{L}}{\partial \hat{y}}$  is propagated back through the network to calculate the error gradient  $d^{(1)}$  at the input layer.

$$\frac{\partial \mathcal{L}}{\partial \hat{y}} \rightarrow d^{(1)} \rightarrow \frac{\partial \mathcal{L}}{\partial z^{(3)}} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} = d^{(1)} \frac{\partial \hat{y}}{\partial z^{(3)}}$$

# Backpropagation: Backward Pass

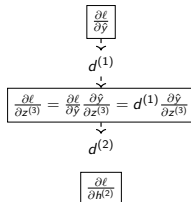
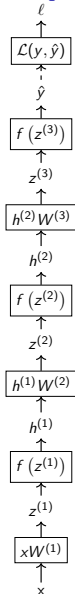




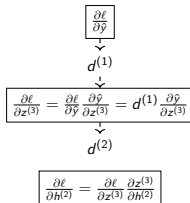
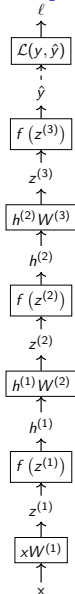
Backward pass diagram showing the flow of error gradients from the loss back to the input:

- The error gradient  $\frac{\partial \mathcal{L}}{\partial \hat{y}}$  is calculated at the loss node.
- This gradient is passed back to the third hidden layer, labeled  $d^{(1)}$ .
- The error gradient is then calculated at the third hidden layer node using the chain rule:  $\frac{\partial \mathcal{L}}{\partial z^{(3)}} = \frac{\partial \mathcal{L}}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z^{(3)}} = d^{(1)} \frac{\partial \hat{y}}{\partial z^{(3)}}$ .
- The resulting gradient is passed back to the second hidden layer, labeled  $d^{(2)}$ .

# Backpropagation: Backward Pass

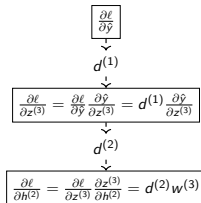
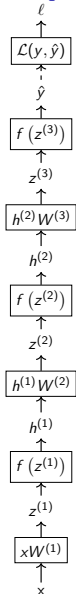


# Backpropagation: Backward Pass

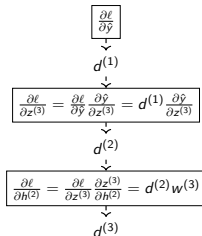
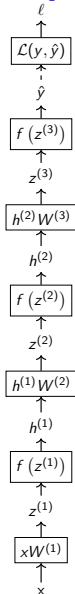




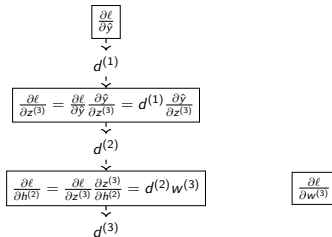
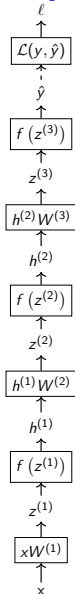
# Backpropagation: Backward Pass



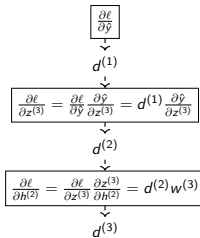
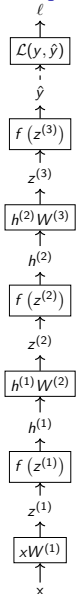
# Backpropagation: Backward Pass



# Backpropagation: Backward Pass

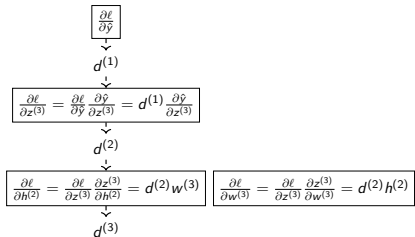
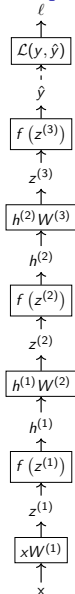


## Backpropagation: Backward Pass

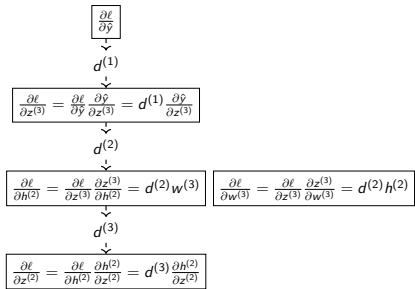
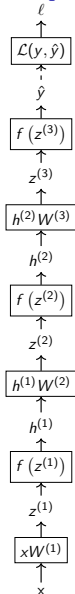


$$\frac{\partial \ell}{\partial \mathbf{w}^{(3)}} = \frac{\partial \ell}{\partial \mathbf{z}^{(3)}} \frac{\partial \mathbf{z}^{(3)}}{\partial \mathbf{w}^{(3)}}$$

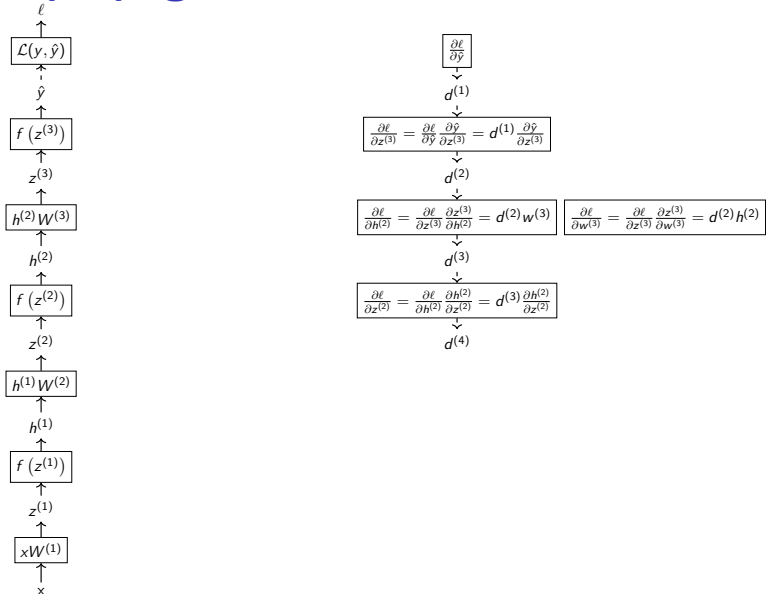
# Backpropagation: Backward Pass



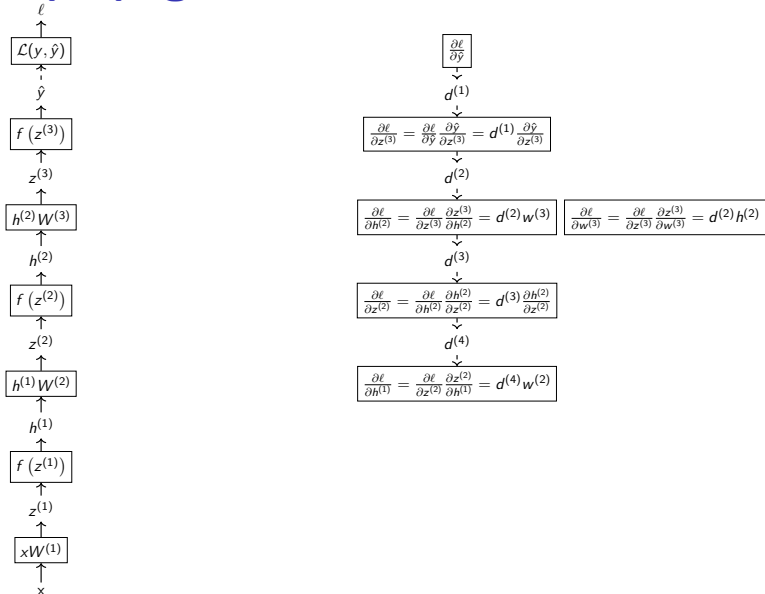
# Backpropagation: Backward Pass



# Backpropagation: Backward Pass

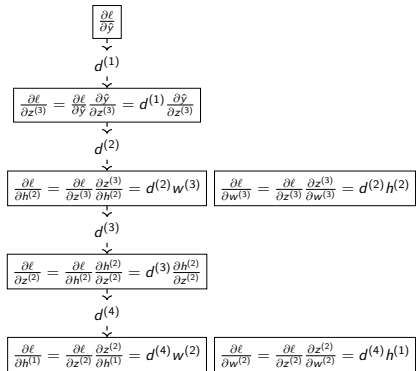
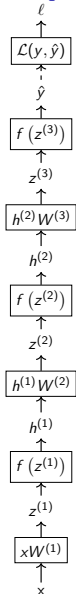


# Backpropagation: Backward Pass

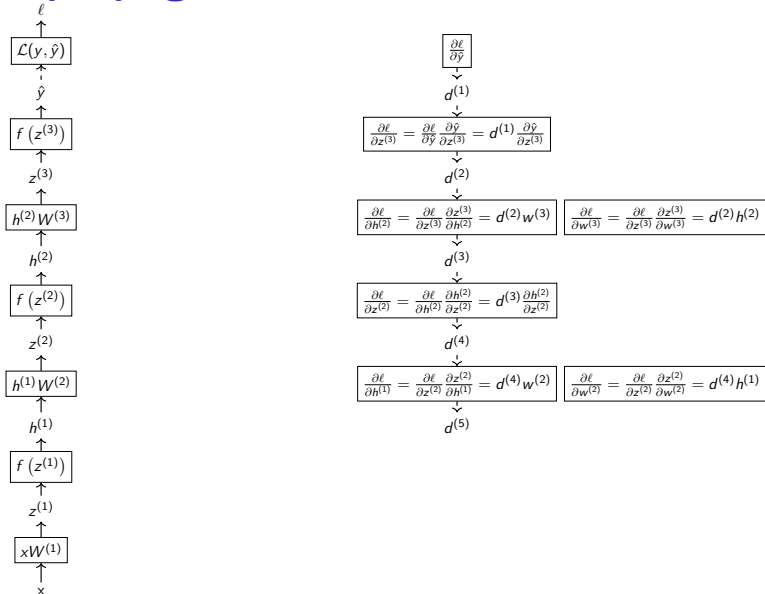




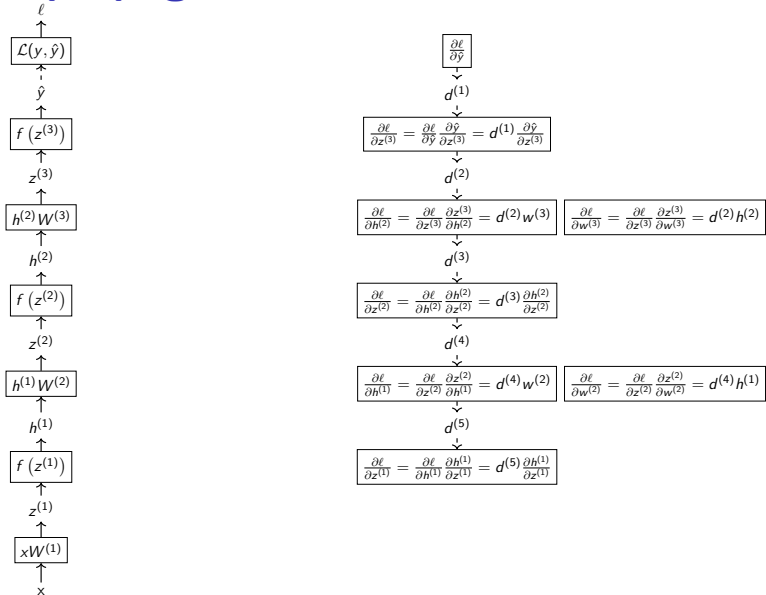
# Backpropagation: Backward Pass



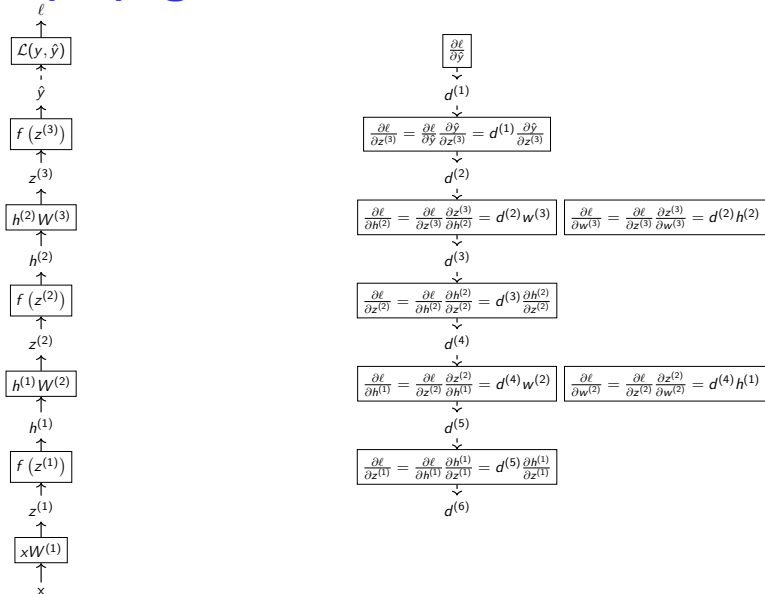
# Backpropagation: Backward Pass



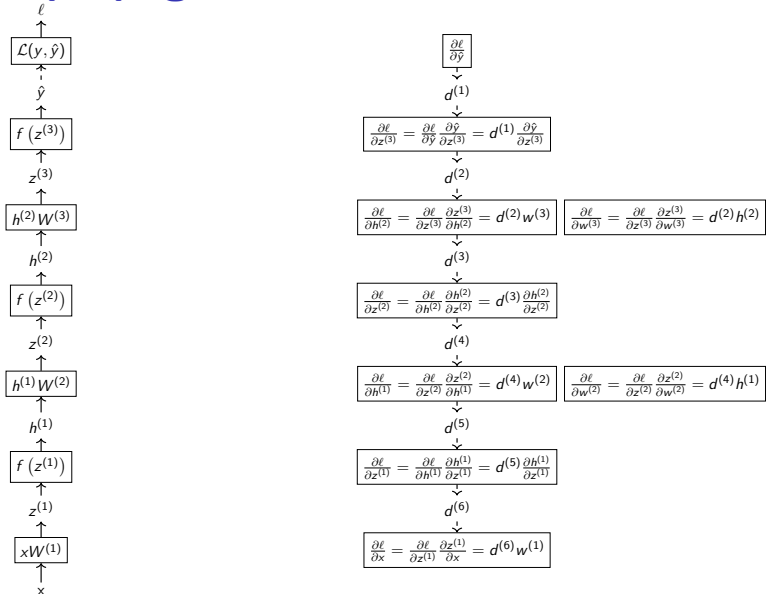
# Backpropagation: Backward Pass



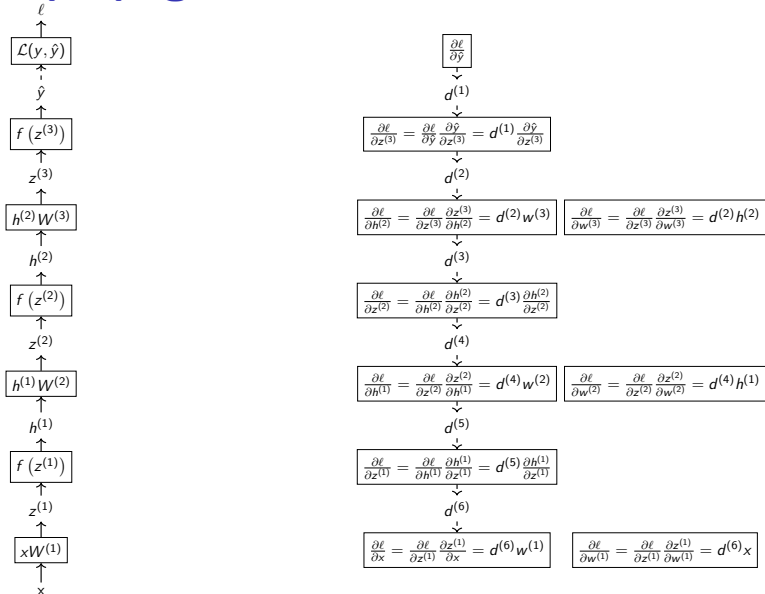
# Backpropagation: Backward Pass



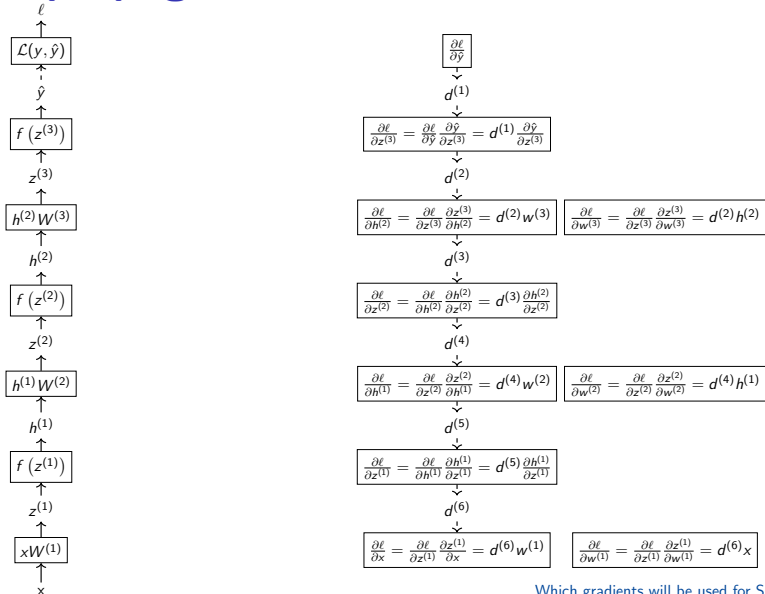
# Backpropagation: Backward Pass



# Backpropagation: Backward Pass

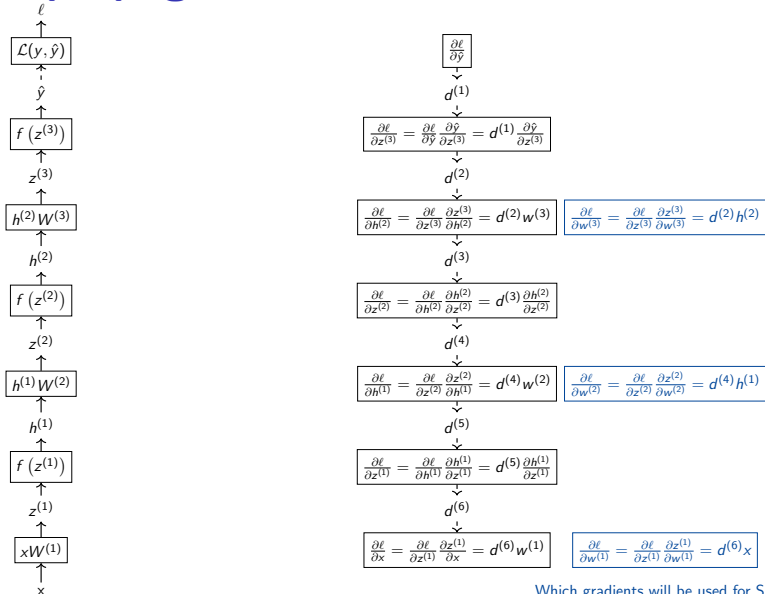


# Backpropagation: Backward Pass



Which gradients will be used for SGD?

# Backpropagation: Backward Pass



Which gradients will be used for SGD?



# Backprop + SGD

- ▶ the output of backprop is gradient of parameters of a neural model

# Backprop + SGD

- ▶ the output of backprop is gradient of parameters of a neural model
- ▶ once we have the gradients we can use SGD rule to update the parameter values

# A Simple Training Loop in PyTorch

```
optimizer = SGD(model_params, lr)

for epoch in range(num_epochs):
    for x,y in data_batches:

        y_hat = model(x)
        loss = loss_func(y_hat, y)

        optimizer.zero_grad()
        loss.backward()

        optimizer.step()
```

# Language Models (LMs)

- ▶ language modeling is the task of assigning a probability to a sentence in a language

# Language Models (LMs)

- ▶ language modeling is the task of assigning a probability to a sentence in a language
- ▶ what is the probability of seeing the sentence “The cat sat on the mat.”

# Language Models (LMs)

- ▶ language modeling is the task of assigning a probability to a sentence in a language
- ▶ what is the probability of seeing the sentence “The cat sat on the mat.”
- ▶ ideal performance at language modeling is to predict the next token in a sequence with a number of guesses that is the identical to or lower than the number of guesses required by a human expert

# Language Models (LMs)

- ▶ language modeling is the task of assigning a probability to a sentence in a language
- ▶ what is the probability of seeing the sentence “The cat sat on the mat.”
- ▶ ideal performance at language modeling is to predict the next token in a sequence with a number of guesses that is the identical to or lower than the number of guesses required by a human expert
- ▶ even without achieving human-level performance, language modeling is a crucial component in real-world NLP applications such as conversational AI, machine-translation, text summarization, ...

# Language Models (LMs)

- ▶ assume a sequence of words  $w_{1:n} = w_1 w_2 \dots w_{n-1} w_n$



# Language Models (LMs)

- ▶ assume a sequence of words  $w_{1:n} = w_1 w_2 \dots w_{n-1} w_n$
- ▶  $P(w_{1:n}) =$   
 $P(w_1)P(w_2|w_1)P(w_3|w_{1:2})P(w_4|w_{1:3})\dots P(w_n|w_{1:n-1})$

# Language Models (LMs)

- ▶ assume a sequence of words  $w_{1:n} = w_1 w_2 \dots w_{n-1} w_n$
- ▶  $P(w_{1:n}) =$   
 $P(w_1)P(w_2|w_1)P(w_3|w_{1:2})P(w_4|w_{1:3})\dots P(w_n|w_{1:n-1})$
- ▶ each word is predicted conditioned on the preceding words

# Language Models (LMs)

- ▶ assume a sequence of words  $w_{1:n} = w_1 w_2 \dots w_{n-1} w_n$
- ▶  $P(w_{1:n}) =$   
 $P(w_1)P(w_2|w_1)P(w_3|w_{1:2})P(w_4|w_{1:3})\dots P(w_n|w_{1:n-1})$
- ▶ each word is predicted conditioned on the preceding words
- ▶ estimating the probability of the last token needs to be conditioned on  $n - 1$  preceding words which is computationally expensive

# Language Models (LMs)

- ▶ assume a sequence of words  $w_{1:n} = w_1 w_2 \dots w_{n-1} w_n$
- ▶  $P(w_{1:n}) =$   
 $P(w_1)P(w_2|w_1)P(w_3|w_{1:2})P(w_4|w_{1:3})\dots P(w_n|w_{1:n-1})$
- ▶ each word is predicted conditioned on the preceding words
- ▶ estimating the probability of the last token needs to be conditioned on  $n - 1$  preceding words which is computationally expensive
- ▶ markov-assumption: the future is independent of the past given the present

# Language Models (LMs)

- ▶  $k$ th order markov-assumption assumes that the next word in a sequence depends only on the last  $k$  words

$$P(w_i | w_{1:i-1}) \approx P(w_i | w_{(i-1)-k:i-1})$$

# Language Models (LMs)

- ▶  $k$ th order markov-assumption assumes that the next word in a sequence depends only on the last  $k$  words

$$P(w_i | w_{1:i-1}) \approx P(w_i | w_{(i-1)-k:i-1})$$

- ▶ probability of a sequence of tokens  $w_{1:n}$

$$P(w_{1:n}) \approx \prod_{i=1}^n P(w_i | w_{i-k:i-1})$$

# Language Models (LMs)

- ▶  $k$ th order markov-assumption assumes that the next word in a sequence depends only on the last  $k$  words

$$P(w_i | w_{1:i-1}) \approx P(w_i | w_{(i-1)-k:i-1})$$

- ▶ probability of a sequence of tokens  $w_{1:n}$

$$P(w_{1:n}) \approx \prod_{i=1}^n P(w_i | w_{i-k:i-1})$$

- ▶ to make it computationally-friendly for computers (What is the problem with the above format?)

$$\log_2 P(w_{1:n}) \approx \sum_{i=1}^n \log_2 (P(w_i | w_{i-k:i-1}))$$

# Evaluating LMs

- ▶ the perplexity metric over an unseen sentence indicates how well a LM predicts the likelihood of the sentence

$$\text{prep}_{w_{1:n}}(\text{LM}) = 2^{-\frac{1}{n} \sum_{i=1}^n \log_2 \text{LM}(w_i | w_{1:i-1})}$$



# Evaluating LMs

- ▶ the perplexity metric over an unseen sentence indicates how well a LM predicts the likelihood of the sentence

$$\text{prep}_{w_{1:n}}(\text{LM}) = 2^{-\frac{1}{n} \sum_{i=1}^n \log_2 \text{LM}(w_i | w_{1:i-1})}$$

- ▶ in this way, we can compare several language models with one another

# Evaluating LMs

- ▶ the perplexity metric over an unseen sentence indicates how well a LM predicts the likelihood of the sentence

$$\text{prep}_{w_{1:n}}(\text{LM}) = 2^{-\frac{1}{n} \sum_{i=1}^n \log_2 \text{LM}(w_i | w_{1:i-1})}$$

- ▶ in this way, we can compare several language models with one another
- ▶ low perplexity values indicate a better language model as it assigns high probabilities to the unseen sentences

# Evaluating LMs

- ▶ the perplexity metric over an unseen sentence indicates how well a LM predicts the likelihood of the sentence

$$\text{prep}_{w_{1:n}}(\text{LM}) = 2^{-\frac{1}{n} \sum_{i=1}^n \log_2 \text{LM}(w_i | w_{1:i-1})}$$

- ▶ in this way, we can compare several language models with one another
- ▶ low perplexity values indicate a better language model as it assigns high probabilities to the unseen sentences
- ▶ perplexities of two language models are only comparable with respect to the same evaluation dataset

# Estimating Probabilities

- ▶ count-based: The estimates can be derived from corpus counts.

# Estimating Probabilities

- ▶ count-based: The estimates can be derived from corpus counts.
- ▶ let  $\#(w_{i:j})$  be the count of the sequence of words  $w_{i:j}$  in a corpus

# Estimating Probabilities

- ▶ count-based: The estimates can be derived from corpus counts.
- ▶ let  $\#(w_{i:j})$  be the count of the sequence of words  $w_{i:j}$  in a corpus
- ▶ the maximum likelihood estimate (MLE) of  $P(w_i|w_{i-1-k:i-1})$

$$P(w_i|w_{i-1-k:i-1}) = \frac{\#(w_{i-1-k:i})}{\#(w_{i-1-k:i-1})}$$

# Estimating Probabilities

- ▶ count-based: The estimates can be derived from corpus counts.
- ▶ let  $\#(w_{i:j})$  be the count of the sequence of words  $w_{i:j}$  in a corpus
- ▶ the maximum likelihood estimate (MLE) of  $P(w_i|w_{i-1-k:i-1})$

$$P(w_i|w_{i-1-k:i-1}) = \frac{\#(w_{i-1-k:i})}{\#(w_{i-1-k:i-1})}$$

- ▶ example:  $w_1 w_2 w_3 = \text{the cat sat}$

$$P(w_3|w_{1:2}) = \frac{\#(\text{the cat sat})}{\#(\text{the cat})}$$

# Estimating Probabilities

- ▶ what if  $\#(w_{i:j}) = 0$ ?



# Estimating Probabilities

- ▶ what if  $\#(w_{i:j}) = 0$ ?
- ▶ then the probability estimation is 0, which translates to an infinite perplexity

# Estimating Probabilities

- ▶ what if  $\#(w_{i:j}) = 0$ ?
- ▶ then the probability estimation is 0, which translates to an infinite perplexity
- ▶ one way of avoiding zero-probability N-grams is to use smoothing techniques

# Estimating Probabilities

- ▶ what if  $\#(w_{i:j}) = 0$ ?
- ▶ then the probability estimation is 0, which translates to an infinite perplexity
- ▶ one way of avoiding zero-probability N-grams is to use smoothing techniques
- ▶ additive smoothing: assume  $|V|$  is the vocabulary size and  $0 < \alpha \leq 1$

$$P(w_i | w_{i-1-k:i-1}) = \frac{\#(w_{i-1-k:i}) + \alpha}{\#(w_{i-1-k:i-1}) + \alpha|V|}$$

# Pro and Cons of Discussed LMs

- ▶ easy to train, scale to large corpora, and work well in practice

# Pro and Cons of Discussed LMs

- ▶ easy to train, scale to large corpora, and work well in practice
- ▶ scaling to larger N-grams is a problem for MLE-based language models.

# Pro and Cons of Discussed LMs

- ▶ easy to train, scale to large corpora, and work well in practice
- ▶ scaling to larger N-grams is a problem for MLE-based language models.
- ▶ the large number of words in the vocabulary means that statistics for larger N-grams will be sparse

# Pro and Cons of Discussed LMs

- ▶ easy to train, scale to large corpora, and work well in practice
- ▶ scaling to larger N-grams is a problem for MLE-based language models.
- ▶ the large number of words in the vocabulary means that statistics for larger N-grams will be sparse
- ▶ MLE-based language models suffer from lack of generalization across contexts

# Pro and Cons of Discussed LMs

- ▶ easy to train, scale to large corpora, and work well in practice
- ▶ scaling to larger N-grams is a problem for MLE-based language models.
- ▶ the large number of words in the vocabulary means that statistics for larger N-grams will be sparse
- ▶ MLE-based language models suffer from lack of generalization across contexts
- ▶ having observed “black car” and “blue car” does not influence our estimates of the sequence “red car” if we haven’t seen it before



# Neural Language Models

- ▶ we can use neural models to estimate probabilities for an LM

# Neural Language Models

- ▶ we can use neural models to estimate probabilities for an LM
- ▶ in this way we can overcome the shortcomings of the MLE-based LMs because neural networks

# Neural Language Models

- ▶ we can use neural models to estimate probabilities for an LM
- ▶ in this way we can overcome the shortcomings of the MLE-based LMs because neural networks
  - ▶ they allow conditioning on increasingly large context sizes with only a linear increase in the number of parameters

# Neural Language Models

- ▶ we can use neural models to estimate probabilities for an LM
- ▶ in this way we can overcome the shortcomings of the MLE-based LMs because neural networks
  - ▶ they allow conditioning on increasingly large context sizes with only a linear increase in the number of parameters
  - ▶ they support generalization across different contexts

# Neural Language Models

- ▶ we can use neural models to estimate probabilities for an LM
- ▶ in this way we can overcome the shortcomings of the MLE-based LMs because neural networks
  - ▶ they allow conditioning on increasingly large context sizes with only a linear increase in the number of parameters
  - ▶ they support generalization across different contexts
- ▶ we focus on the neural LM that was introduced by Bengio et al. (2003)

# Neural Language Models

- ▶ let  $w_{1:k}$  be the given context

# Neural Language Models

- ▶ let  $w_{1:k}$  be the given context
- ▶ we want to estimate  $P(w_{k+1}|w_{1:k})$

# Neural Language Models

- ▶ let  $w_{1:k}$  be the given context
- ▶ we want to estimate  $P(w_{k+1}|w_{1:k})$
- ▶ we design an MLP neural model, which takes  $w_{1:k}$  as input and returns  $P(w_{k+1})$  over all words in vocabulary  $V$  as output

$$x = [v(w_1), v(w_2), \dots, v(w_k)]$$

$$h^{(1)} = g(xW^{(1)} + b^{(1)})$$

$$P(w_{k+1}) = \text{softmax}(h^{(1)}W^{(2)} + b^{(2)})$$



# Neural Language Models

- ▶ let  $w_{1:k}$  be the given context
- ▶ we want to estimate  $P(w_{k+1}|w_{1:k})$
- ▶ we design an MLP neural model, which takes  $w_{1:k}$  as input and returns  $P(w_{k+1})$  over all words in vocabulary  $V$  as output

$$x = [v(w_1), v(w_2), \dots, v(w_k)]$$

$$h^{(1)} = g(xW^{(1)} + b^{(1)})$$

$$P(w_{k+1}) = \text{softmax}(h^{(1)}W^{(2)} + b^{(2)})$$

- ▶ the training examples are simply word  $k$ -grams from the training set, where the identities of the first  $k - 1$  words are used as features, and the last word is used as the target label for the classification

# Neural Language Models

- ▶ let  $w_{1:k}$  be the given context
- ▶ we want to estimate  $P(w_{k+1}|w_{1:k})$
- ▶ we design an MLP neural model, which takes  $w_{1:k}$  as input and returns  $P(w_{k+1})$  over all words in vocabulary  $V$  as output

$$x = [v(w_1), v(w_2), \dots, v(w_k)]$$

$$h^{(1)} = g(xW^{(1)} + b^{(1)})$$

$$P(w_{k+1}) = \text{softmax}(h^{(1)}W^{(2)} + b^{(2)})$$

- ▶ the training examples are simply word  $k$ -grams from the training set, where the identities of the first  $k - 1$  words are used as features, and the last word is used as the target label for the classification
- ▶ loss function: cross-entropy loss

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context
- ▶ we are asked to predict the next word  $w_{k+1}$  from vocabulary  $V$

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context
- ▶ we are asked to predict the next word  $w_{k+1}$  from vocabulary  $V$
- ▶ we feed  $w_{1:k}$  to our trained MLP-based LM

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context
- ▶ we are asked to predict the next word  $w_{k+1}$  from vocabulary  $V$
- ▶ we feed  $w_{1:k}$  to our trained MLP-based LM
- ▶ our LM returns  $P(w_{k+1})$  of each word in  $V$

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context
- ▶ we are asked to predict the next word  $w_{k+1}$  from vocabulary  $V$
- ▶ we feed  $w_{1:k}$  to our trained MLP-based LM
- ▶ our LM returns  $P(w_{k+1})$  of each word in  $V$
- ▶ we pick up the word with the maximum probability to generate the next word

# Neural LMs for Generating Language

- ▶ assume that we are given  $w_{1:k}$  as context
- ▶ we are asked to predict the next word  $w_{k+1}$  from vocabulary  $V$
- ▶ we feed  $w_{1:k}$  to our trained MLP-based LM
- ▶ our LM returns  $P(w_{k+1})$  of each word in  $V$
- ▶ we pick up the word with the maximum probability to generate the next word
- ▶ we add the the predicted word to the context and repeat the above procedure



# Summary

- ▶ training as optimization of a loss function
- ▶ common loss functions and regularization terms
- ▶ gradient descent (GD, SGD): A general technique for optimization
- ▶ backprop(agation): An algorithm for deriving gradients in neural models, once gradients are determined we can train a model with SGD
- ▶ (Neural) Language Models (LMs)

Thank You!