

# Deep Learning for NLP

## Lecture 4: Text Representations I

**Dr. Mohsen Mesgar**

**Ubiquitous Knowledge Processing Lab (UKP Lab)**

# This lecture

- ▶ common features used for converting textual data into numerical vectors
- ▶ basics of word embeddings
  - ▶ how to get them?
  - ▶ where to use them?
  - ▶ how to use them?
- ▶ limitations

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples

# Recall

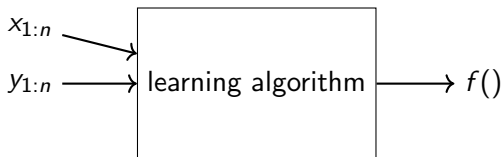
- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels

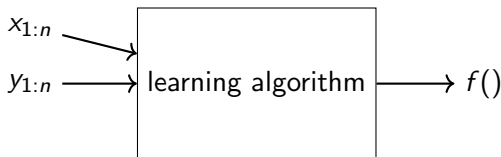
# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels



# Recall

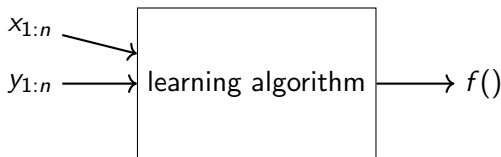
- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels





# Recall

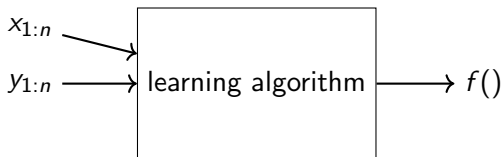
- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels



- ▶ the input to neural models  $f()$  should be a numerical vector but NLP tasks are defined on texts

# Recall

- ▶ the input to a supervised learning algorithm is a training set  $(x_{1:n}, y_{1:n})$ , where
  - ▶  $x_{1:n} = x_1, x_2, \dots, x_n$  shows input examples
  - ▶  $y_{1:n} = y_1, y_2, \dots, y_n$  shows corresponding labels
- ▶ the goal of a learning algorithm is to return a function  $f()$  that accurately maps input examples to their desired labels




- ▶ the input to neural models  $f()$  should be a numerical vector but NLP tasks are defined on texts
- ▶ how can we represent texts via numerical vectors?

# Text Representation (a.k.a., Feature Extraction)

- ▶ vector representation of a text should ideally reflect various linguistic properties of the text

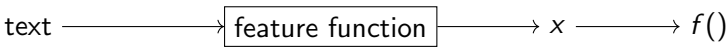
# Text Representation (a.k.a., Feature Extraction)

- ▶ vector representation of a text should ideally reflect various linguistic properties of the text



# Text Representation (a.k.a., Feature Extraction)

- ▶ vector representation of a text should ideally reflect various linguistic properties of the text

- ▶ 

```
graph LR; text --> ff[feature function]; ff --> x; x --> f["f()"]
```
- ▶ in this lecture we focus on feature functions rather than learning algorithms

# Some Terminologies

- ▶ letters: smallest units in a language  $\rightarrow$  a,b,c,...

# Some Terminologies

- ▶ letters: smallest units in a language  $\rightarrow$  a,b,c,...
- ▶ tokens and words: tokens are outputs of a tokenizer and words are meaning-bearing units
  - ▶ note: in this course we use the terms “word” and “token” interchangeably

# Some Terminologies

- ▶ letters: smallest units in a language → a,b,c,...
- ▶ tokens and words: tokens are outputs of a tokenizer and words are meaning-bearing units
  - ▶ note: in this course we use the terms “word” and “token” interchangeably
- ▶ lemma: the dictionary entry of the word → “book” is the lemma of “booking”, “booked”, and “books”
  - ▶ how to obtain lemmas? use morphological analyzers
  - ▶ are available for many languages
  - ▶ lemmatization may not work for any sequence of letters e.g., for mis-spelling



# Some Terminologies

- ▶ letters: smallest units in a language → a,b,c,...
- ▶ tokens and words: tokens are outputs of a tokenizer and words are meaning-bearing units
  - ▶ note: in this course we use the terms “word” and “token” interchangeably
- ▶ lemma: the dictionary entry of the word → “book” is the lemma of “booking”, “booked”, and “books”
  - ▶ how to obtain lemmas? use morphological analyzers
  - ▶ are available for many languages
  - ▶ lemmatization may not work for any sequence of letters e.g., for mis-spelling
- ▶ stems: a shorter venison of a word defined based on some language-specific heuristic → “pictur” is the stem of “pictures”, “pictured”, and “picture”
  - ▶ the output of a stemmer need not be a valid word

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs



# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English
    - ▶ focus around verbs

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English
    - ▶ focus around verbs
    - ▶ manually created

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English
    - ▶ focus around verbs
    - ▶ manually created
  - ▶ Paraphrase Database (PPDB)

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English
    - ▶ focus around verbs
    - ▶ manually created
  - ▶ Paraphrase Database (PPDB)
    - ▶ contains paraphrases

# Some Terminologies

- ▶ lexical resources: provide some information about words and their relations
  - ▶ dictionaries
  - ▶ WordNet for English
    - ▶ capture semantic knowledge about words
    - ▶ each word is associated with one or several synsets
    - ▶ synsets are linked to each other according to semantic relations between their words
    - ▶ semantic relations are: hypernym (more specific), hyponym (less specific), antonyms, holonyms (part-whole), and meronyms (whole-part)
    - ▶ contains information about nouns, verbs, adjectives, and adverbs
    - ▶ manually created
  - ▶ FrameNet and VerbNet for English
    - ▶ focus around verbs
    - ▶ manually created
  - ▶ Paraphrase Database (PPDB)
    - ▶ contains paraphrases
    - ▶ automatically created

# Common Features for Textual Data

what information can we extract directly from a word

- ▶ letters comprising a word and their order



# Common Features for Textual Data

what information can we extract directly from a word

- ▶ letters comprising a word and their order
- ▶ length of word

# Common Features for Textual Data

what information can we extract directly from a word

- ▶ letters comprising a word and their order
- ▶ length of word
- ▶ is the first letter capitalized?

# Common Features for Textual Data

what information can we extract directly from a word

- ▶ letters comprising a word and their order
- ▶ length of word
- ▶ is the first letter capitalized?
- ▶ does word include hyphen?


# Common Features for Textual Data

- Bag-of-Words (BoW): the count of each word in a text is taken as a feature

## The Bag of Words Representation

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!

15



it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
would	1
whimsical	1
times	1
sweet	1
satirical	1
adventure	1
genre	1
fairy	1
humor	1
have	1
great	1
...	...

(Taken from: <https://www.programmersonsought.com/article/4304366575/>)

# Common Features for Textual Data

- Bag-of-Words (BoW): the count of each word in a text is taken as a feature

## The Bag of Words Representation

I love this movie! It's sweet, but with satirical humor. The dialogue is great and the adventure scenes are fun... It manages to be whimsical and romantic while laughing at the conventions of the fairy tale genre. I would recommend it to just about anyone. I've seen it several times, and I'm always happy to see it again whenever I have a friend who hasn't seen it yet!

15



it	6
I	5
the	4
to	3
and	3
seen	2
yet	1
would	1
whimsical	1
times	1
sweet	1
satirical	1
adventure	1
genre	1
fairy	1
humor	1
have	1
great	1
...	...

- BOW does not care about the order of words

(Taken from: <https://www.programmingsought.com/article/4304366575/>)

# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency

# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency
- ▶ let  $d$  be a document from a given corpus  $D$

# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency
- ▶ let  $d$  be a document from a given corpus  $D$
- ▶ we map word  $w$  of  $d$  to a number as follows:

$$\frac{\#(w, d)}{\sum_{w'} \#(w', d)} \times \log \frac{|D|}{|\{d \in D : w \in d\}|}$$



# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency
- ▶ let  $d$  be a document from a given corpus  $D$
- ▶ we map word  $w$  of  $d$  to a number as follows:

$$\frac{\#(w, d)}{\sum_{w'} \#(w', d)} \times \log \frac{|D|}{|\{d \in D : w \in d\}|}$$

- ▶ N-grams: instead of using the frequency of a word, we use the frequency of N sequence of words

# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency
- ▶ let  $d$  be a document from a given corpus  $D$
- ▶ we map word  $w$  of  $d$  to a number as follows:

$$\frac{\#(w, d)}{\sum_{w'} \#(w', d)} \times \log \frac{|D|}{|\{d \in D : w \in d\}|}$$

- ▶ N-grams: instead of using the frequency of a word, we use the frequency of N sequence of words
- ▶  $N=2 \rightarrow$  bi-gram

# Common Features for Textual Data

- ▶ TF-IDF: Term Frequency - Inverse Document Frequency
- ▶ let  $d$  be a document from a given corpus  $D$
- ▶ we map word  $w$  of  $d$  to a number as follows:

$$\frac{\#(w, d)}{\sum_{w'} \#(w', d)} \times \log \frac{|D|}{|\{d \in D : w \in d\}|}$$

- ▶ N-grams: instead of using the frequency of a word, we use the frequency of N sequence of words
- ▶  $N=2 \rightarrow$  bi-gram
- ▶  $N=3 \rightarrow$  tri-gram

# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context

# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context
- ▶ how to encode contextual information?

# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context
- ▶ how to encode contextual information?
- ▶ using position of a word within a sentence or a document

# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context
- ▶ how to encode contextual information?
- ▶ using position of a word within a sentence or a document
- ▶ using the words that appear in an immediate context of a word

# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context
- ▶ how to encode contextual information?
- ▶ using position of a word within a sentence or a document
- ▶ using the words that appear in an immediate context of a word
  - ▶ immediate context: a window of  $k$  words surrounding the target word



# Common Features for Textual Data

- ▶ context: each piece of a text occurs within a larger text which is known as context
- ▶ how to encode contextual information?
- ▶ using position of a word within a sentence or a document
- ▶ using the words that appear in an immediate context of a word
  - ▶ immediate context: a window of  $k$  words surrounding the target word
  - ▶ “the cat sat on the mat” ,  $k=2 \rightarrow \{ \text{word-minus-2=the, word-minus-1=cat, word-plus-1=on, word-plus-2=the} \}$

# Common Features for Textual Data

what information can we extract from the relation of text with external source of information?

- ▶ is a word in a list of common person names in Germany?

# Common Features for Textual Data

what information can we extract from the relation of text with external source of information?

- ▶ is a word in a list of common person names in Germany?
- ▶ is a word a female or male person name?

# Common Features for Textual Data

what information can we extract from the relation of text with external source of information?

- ▶ is a word in a list of common person names in Germany?
- ▶ is a word a female or male person name?
- ▶ what is the lemma of the word?

# Common Features for Textual Data

what information can we extract from the relation of text with external source of information?

- ▶ is a word in a list of common person names in Germany?
- ▶ is a word a female or male person name?
- ▶ what is the lemma of the word?
- ▶ what is the stem of the word?

# Common Features for Textual Data

what information can we extract from the relation of text with external source of information?

- ▶ is a word in a list of common person names in Germany?
- ▶ is a word a female or male person name?
- ▶ what is the lemma of the word?
- ▶ what is the stem of the word?
- ▶ what information do lexical resources give us about the word?

# Feature Vectors

- ▶ the output of a learning algorithm is a  $f()$

# Feature Vectors

- ▶ the output of a learning algorithm is a  $f()$
- ▶  $f$  takes as input a vector  $x$  with dimension  $d_{in}$



# Feature Vectors

- ▶ the output of a learning algorithm is a  $f()$
- ▶  $f$  takes as input a vector  $x$  with dimension  $d_{in}$
- ▶  $f$  returns as output a vector  $\hat{y}$  with dimension  $d_{out}$

# Feature Vectors

- ▶ the output of a learning algorithm is a  $f()$
- ▶  $f$  takes as input a vector  $x$  with dimension  $d_{in}$
- ▶  $f$  returns as output a vector  $\hat{y}$  with dimension  $d_{out}$
- ▶ how can we map textual features to a vector?

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text
- ▶ categorical features
  - ▶ the value of a feature is from a set of values

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text
- ▶ categorical features
  - ▶ the value of a feature is from a set of values
  - ▶ what is the POS of a word?

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text
- ▶ categorical features
  - ▶ the value of a feature is from a set of values
  - ▶ what is the POS of a word?
- ▶ how to encode categorical features?

# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text
- ▶ categorical features
  - ▶ the value of a feature is from a set of values
  - ▶ what is the POS of a word?
- ▶ how to encode categorical features?
  - ▶ one-hot encodings



# Different Types of Features

- ▶ numerical features
  - ▶ the value of a feature is a number
  - ▶ e.g., the frequency of a word in a text
- ▶ categorical features
  - ▶ the value of a feature is from a set of values
  - ▶ what is the POS of a word?
- ▶ how to encode categorical features?
  - ▶ one-hot encodings
  - ▶ dense embeddings

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$
- ▶ we encode the space of feature values via vectors in which

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$
- ▶ we encode the space of feature values via vectors in which
  - ▶ each entry is either 0 or 1

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$
- ▶ we encode the space of feature values via vectors in which
  - ▶ each entry is either 0 or 1
  - ▶ each entry is associated with one of the categories

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$
- ▶ we encode the space of feature values via vectors in which
  - ▶ each entry is either 0 or 1
  - ▶ each entry is associated with one of the categories
  - ▶ only one item in a vector can be 1, the rest should be 0

# One-hot Encodings

- ▶ let's assume that values of a feature is in categories  $\{v_0, v_1, v_2\}$
- ▶ we encode the space of feature values via vectors in which
  - ▶ each entry is either 0 or 1
  - ▶ each entry is associated with one of the categories
  - ▶ only one item in a vector can be 1, the rest should be 0
- ▶ for example:
  - ▶  $v_0 = [1, 0, 0]$
  - ▶  $v_1 = [0, 1, 0]$
  - ▶  $v_2 = [0, 0, 1]$

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$



# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?
  - ▶  $\text{news} = [1, 0, 0, 0]$

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?
  - ▶ news =  $[1, 0, 0, 0]$
  - ▶ sport =  $[0, 1, 0, 0]$

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?
  - ▶ news =  $[1, 0, 0, 0]$
  - ▶ sport =  $[0, 1, 0, 0]$
  - ▶ science =  $[0, 0, 1, 0]$

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?
  - ▶ news =  $[1, 0, 0, 0]$
  - ▶ sport =  $[0, 1, 0, 0]$
  - ▶ science =  $[0, 0, 1, 0]$
  - ▶ politic =  $[0, 0, 0, 1]$

# Example: Categorical Labels

- ▶ let  $f()$  take a vector  $x$  which encodes a text and also return a vector  $\hat{y}$  representing the topic of the text which can be from  $\{\text{news, sport, science, politic}\}$
- ▶ how can we represent topic labels using one-hot encoding?
  - ▶ news =  $[1, 0, 0, 0]$
  - ▶ sport =  $[0, 1, 0, 0]$
  - ▶ science =  $[0, 0, 1, 0]$
  - ▶ politic =  $[0, 0, 0, 1]$
- ▶ what is the length of one-hot vectors for a feature with  $k$  categories?

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text



# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$ 
  - ▶  $v_0 = [1, 0, 0, \dots, 0]$

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$ 
  - ▶  $v_0 = [1, 0, 0, \dots, 0]$
  - ▶  $v_1 = [0, 1, 0, \dots, 0]$

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$ 
  - ▶  $v_0 = [1, 0, 0, \dots, 0]$
  - ▶  $v_1 = [0, 1, 0, \dots, 0]$
  - ▶  $v_2 = [0, 0, 1, \dots, 0]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [0, 0, 0, \dots, 1]$

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$ 
  - ▶  $v_0 = [1, 0, 0, \dots, 0]$
  - ▶  $v_1 = [0, 1, 0, \dots, 0]$
  - ▶  $v_2 = [0, 0, 1, \dots, 0]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [0, 0, 0, \dots, 1]$
- ▶ we can easily represent a text via BOW and then represent each word in BOW with its one-hot vector

# Example: Word Encodings

- ▶ let  $V$  be vocabulary of a language
- ▶ we take  $V$  as categories a word can take in a text
- ▶ so words can be encoded by one-hot vectors
- ▶ let  $V = \{v_0, v_1, v_2, \dots, v_{|V|-1}\}$ 
  - ▶  $v_0 = [1, 0, 0, \dots, 0]$
  - ▶  $v_1 = [0, 1, 0, \dots, 0]$
  - ▶  $v_2 = [0, 0, 1, \dots, 0]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [0, 0, 0, \dots, 1]$
- ▶ we can easily represent a text via BOW and then represent each word in BOW with its one-hot vector
- ▶ how can we define vocabulary  $V$  given a corpus  $D$ ?

# Comment

pause



# One-hot Encodings

- ▶ word vectors are very sparse

# One-hot Encodings

- ▶ word vectors are very sparse
- ▶ semantic relations between words are not encoded in word vectors

# One-hot Encodings

- ▶ word vectors are very sparse
- ▶ semantic relations between words are not encoded in word vectors
- ▶ it's better to use one-hot representations for a few distinct features where we expect no correlation between features

# Dense Encodings

- ▶ a categorical feature is embedded as a vector in a  $d$  dimensional space

# Dense Encodings

- ▶ a categorical feature is embedded as a vector in a  $d$  dimensional space
- ▶ assuming categories  $C = \{c_0, c_1, \dots, c_{|C|-1}\}$ ,  $d = 4$

# Dense Encodings

- ▶ a categorical feature is embedded as a vector in a  $d$  dimensional space
- ▶ assuming categories  $C = \{c_0, c_1, \dots, c_{|C|-1}\}$ ,  $d = 4$ 
  - ▶  $c_0 = [+0.1, -0.2, +0.3, +0.5]$

# Dense Encodings

- ▶ a categorical feature is embedded as a vector in a  $d$  dimensional space
- ▶ assuming categories  $C = \{c_0, c_1, \dots, c_{|C|-1}\}$ ,  $d = 4$ 
  - ▶  $c_0 = [+0.1, -0.2, +0.3, +0.5]$
  - ▶  $c_1 = [-0.2 - 0.1, +0.1, +0.2]$
  - ▶ ...
  - ▶  $c_{|C|-1} = [+0.2, -0.2, -0.1, +0.3]$

# Example: Word Encodings

- ▶ assuming vocabulary  $V = \{v_0, v_1, \dots, v_{|V|-1}\}$ ,  $d = 4$



# Example: Word Encodings

- ▶ assuming vocabulary  $V = \{v_0, v_1, \dots, v_{|V|-1}\}$ ,  $d = 4$ 
  - ▶  $v_0 = [+0.1, -0.2, +0.3, +0.5]$
  - ▶  $v_1 = [-0.2, -0.1, +0.1, +0.2]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [+0.2, -0.2, -0.1, +0.3]$

# Example: Word Encodings

- ▶ assuming vocabulary  $V = \{v_0, v_1, \dots, v_{|V|-1}\}$ ,  $d = 4$ 
  - ▶  $v_0 = [+0.1, -0.2, +0.3, +0.5]$
  - ▶  $v_1 = [-0.2, -0.1, +0.1, +0.2]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [+0.2, -0.2, -0.1, +0.3]$
- ▶ in the one-hot method we represent each word with a sparse vector of size  $|V|$

# Example: Word Encodings

- ▶ assuming vocabulary  $V = \{v_0, v_1, \dots, v_{|V|-1}\}$ ,  $d = 4$ 
  - ▶  $v_0 = [+0.1, -0.2, +0.3, +0.5]$
  - ▶  $v_1 = [-0.2, -0.1, +0.1, +0.2]$
  - ▶ ...
  - ▶  $v_{|V|-1} = [+0.2, -0.2, -0.1, +0.3]$
- ▶ in the one-hot method we represent each word with a sparse vector of size  $|V|$
- ▶ in dense encoding method we represent each word with a dense vector with a small size  $d$

# Dense Encodings

- ▶ dimensionality of vectors is  $d$

# Dense Encodings

- ▶ dimensionality of vectors is  $d$
- ▶ model training will cause similar features to have similar vectors

# Dense Encodings

- ▶ dimensionality of vectors is  $d$
- ▶ model training will cause similar features to have similar vectors
- ▶ it's mainly useful when we expect some correlations between features → “cat” and “dog” are semantically related

# Dense Encodings

- ▶ dimensionality of vectors is  $d$
- ▶ model training will cause similar features to have similar vectors
- ▶ it's mainly useful when we expect some correlations between features → “cat” and “dog” are semantically related
- ▶ is also useful when we have a large number of features → for example vocabulary

- ▶ by embeddings we mean representing each textual feature as a dense vector in a low dimensional space



- ▶ by embeddings we mean representing each textual feature as a dense vector in a low dimensional space
- ▶ how should we define these dense vectors or embeddings?

# Word Embeddings

- ▶ dense representations of words in an embedding space is known as word embeddings
- ▶ how to obtain word embeddings?

# Random Initialization

- ▶ we initialize the embedding vectors to random values

# Random Initialization

- ▶ we initialize the embedding vectors to random values
- ▶ values are uniformly sampled from numbers in the range

# Random Initialization

- ▶ we initialize the embedding vectors to random values
- ▶ values are uniformly sampled from numbers in the range
  - ▶  $[-\frac{1}{2d}, +\frac{1}{2d}]$  where  $d$  is the number of dimensions (Mikolov et al., 2013)

# Random Initialization

- ▶ we initialize the embedding vectors to random values
- ▶ values are uniformly sampled from numbers in the range
  - ▶  $[-\frac{1}{2d}, +\frac{1}{2d}]$  where  $d$  is the number of dimensions (Mikolov et al., 2013)
  - ▶  $[-\frac{\sqrt{6}}{\sqrt{d}}, +\frac{\sqrt{6}}{\sqrt{d}}]$  (xavier initialization)

# Word Embeddings

- ▶ dense representations of words in an embedding space is known as word embeddings

# Word Embeddings

- ▶ dense representations of words in an embedding space is known as word embeddings
- ▶ these vectors can be obtained by random initialization and tuned for any NLP task



# Word Embeddings

- ▶ dense representations of words in an embedding space is known as word embeddings
- ▶ these vectors can be obtained by random initialization and tuned for any NLP task
- ▶ however, we as human beings can define semantic relations between words independent of any NLP task
  - ▶ e.g., “blue” and “black” are colors
  - ▶ e.g., “dog” is more similar to “cat” than to “chair”
  - ▶ e.g., “easy” is the opposite of “difficult”

# Word Embeddings

- ▶ dense representations of words in an embedding space is known as word embeddings
- ▶ these vectors can be obtained by random initialization and tuned for any NLP task
- ▶ however, we as human beings can define semantic relations between words independent of any NLP task
  - ▶ e.g., “blue” and “black” are colors
  - ▶ e.g., “dog” is more similar to “cat” than to “chair”
  - ▶ e.g., “easy” is the opposite of “difficult”
- ▶ how can we find word embeddings such that vectors of words with similar meaning be close to each other in the embedding space?

# Word Meaning

- ▶ “you should know a word by the company it keeps” (Frith, 1957)



# Distributional Hypothesis

- ▶ words that occur in the same contexts tend to have similar meanings (Harris, 1945)

# Distributional Hypothesis

- ▶ words that occur in the same contexts tend to have similar meanings (Harris, 1945)
- ▶ how can we model the distributional hypothesis?

# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus

# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus
- ▶ context is defined by a window over words

# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus
- ▶ context is defined by a window over words
- ▶ let  $V$  be the set of words in vocabulary and  $C$  be the set of possible contexts



# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus
- ▶ context is defined by a window over words
- ▶ let  $V$  be the set of words in vocabulary and  $C$  be the set of possible contexts
- ▶ word-context matrix  $M$  is a two dimensional matrix whose rows are associated with  $V$  and columns with  $C$

# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus
- ▶ context is defined by a window over words
- ▶ let  $V$  be the set of words in vocabulary and  $C$  be the set of possible contexts
- ▶ word-context matrix  $M$  is a two dimensional matrix whose rows are associated with  $V$  and columns with  $C$
- ▶ each entry of the matrix indicates how often a word co-occurs with a context in the given corpus

# Word-Context Matrix

- ▶ we count how often a word has occurred with a context in texts from a large corpus
- ▶ context is defined by a window over words
- ▶ let  $V$  be the set of words in vocabulary and  $C$  be the set of possible contexts
- ▶ word-context matrix  $M$  is a two dimensional matrix whose rows are associated with  $V$  and columns with  $C$
- ▶ each entry of the matrix indicates how often a word co-occurs with a context in the given corpus
- ▶ this matrix is also known as co-occurrence matrix

# Word-Context Matrix

- ▶ corpus:
  - ▶ I like DL
  - ▶ I like NLP
  - ▶ I love ML
  - ▶ I love NLP
- ▶ window size = 1

# Word-Context Matrix

- ▶ corpus:
  - ▶ I like DL
  - ▶ I like NLP
  - ▶ I love ML
  - ▶ I love NLP
- ▶ window size = 1

	I	like	love	DL	NLP	ML
I	0	2	2	0	0	0
like	2	0	0	1	1	0
love	2	0	0	0	1	1
DL	0	1	0	0	0	0
NLP	0	1	1	0	0	0
ML	0	0	1	0	0	0

# Word-Context Matrix

- ▶ corpus:
  - ▶ I like DL
  - ▶ I like NLP
  - ▶ I love ML
  - ▶ I love NLP
- ▶ window size = 1

	I	like	love	DL	NLP	ML
I	0	2	2	0	0	0
like	2	0	0	1	1	0
love	2	0	0	0	1	1
DL	0	1	0	0	0	0
NLP	0	1	1	0	0	0
ML	0	0	1	0	0	0

# Count-based Embeddings

- ▶ for a large corpus, the size of matrix is very large

# Count-based Embeddings

- ▶ for a large corpus, the size of matrix is very large
- ▶ the matrix could be sparse too



# Count-based Embeddings

- ▶ for a large corpus, the size of matrix is very large
- ▶ the matrix could be sparse too
- ▶ to encounter this, we may use dimensionality reduction techniques

# Count-based Embeddings

- ▶ for a large corpus, the size of matrix is very large
- ▶ the matrix could be sparse too
- ▶ to encounter this, we may use dimensionality reduction techniques
- ▶ however, for adding a new word we need to enlarge the matrix and apply dimensionality reduction again

# CBoW Method

- ▶ CBoW stands for Continuous Bag of Words
- ▶ task: given a context  $\rightarrow$  predict a missing word from the context

input: "I \_ NLP"  $\rightarrow$  output: "like"

# CBoW Method

- ▶ CBoW stands for Continuous Bag of Words
- ▶ task: given a context  $\rightarrow$  predict a missing word from the context

input: "I \_ NLP"  $\rightarrow$  output: "like"

input: ("I \_ NLP", "like")  $\rightarrow$  output: 1 and

input: ("I \_ NLP", "apple")  $\rightarrow$  output: 0

# CBoW Method

- ▶ given a corpus, we create two sets:

# CBoW Method

- ▶ given a corpus, we create two sets:
  - ▶ positive set ( $D$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is the correct value for the missing word

# CBoW Method

- ▶ given a corpus, we create two sets:
  - ▶ positive set ( $D$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is the correct value for the missing word
  - ▶ negative set ( $D'$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is a random value for the missing word

# CBoW Method

- ▶ given a corpus, we create two sets:
  - ▶ positive set ( $D$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is the correct value for the missing word
  - ▶ negative set ( $D'$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is a random value for the missing word
- ▶ we compute the score estimating similarity between context  $c$  and word  $w$  given context-word pair  $(c, w)$

$$s(c, w) = e(w) \sum_{w_i \in c} e(w_i)$$

where  $e$  is a function that maps each word to its embeddings



# CBoW Method

- ▶ given a corpus, we create two sets:
  - ▶ positive set ( $D$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is the correct value for the missing word
  - ▶ negative set ( $D'$ ): consisting of pairs  $(c, w)$  where  $c$  is a context and  $w$  is a random value for the missing word
- ▶ we compute the score estimating similarity between context  $c$  and word  $w$  given context-word pair  $(c, w)$

$$s(c, w) = e(w) \sum_{w_i \in c} e(w_i)$$

where  $e$  is a function that maps each word to its embeddings

- ▶ example

$$s(\text{I} \_ \text{NLP}, \text{like}) = e(\text{like})(e(\text{I}) + e(\text{NLP}))$$

# CBoW Method

- ▶ we use the sigmoid function to map the score to a probability

$$P(y = 1|(c, w)) = \frac{1}{1 + e^{-s(c, w)}}$$

# CBoW Method

- ▶ we use the sigmoid function to map the score to a probability

$$P(y = 1|(c, w)) = \frac{1}{1 + e^{-s(c, w)}}$$

- ▶ we should maximize:

$$\frac{1}{|D|} \sum_{(c, w) \in D} P(y = 1|(c, w)) + \frac{1}{|D'|} \sum_{(c, w) \in D'} P(y = 0|(c, w))$$

# CBoW Method

- ▶ we use the sigmoid function to map the score to a probability

$$P(y = 1|(c, w)) = \frac{1}{1 + e^{-s(c, w)}}$$

- ▶ we should maximize:

$$\frac{1}{|D|} \sum_{(c, w) \in D} P(y = 1|(c, w)) + \frac{1}{|D'|} \sum_{(c, w) \in D'} P(y = 0|(c, w))$$

or minimize the following loss

$$L(\Theta) = -\frac{1}{|D|} \sum_{(c, w) \in D} \log P(y = 1|(c, w)) - \frac{1}{|D'|} \sum_{(c, w) \in D'} \log P(y = 0|(c, w))$$

# Skip-Gram Method

- ▶ we treat words of a context independent from each other

$$P(y = 1|(c, w)) = \prod_{c_i \in c} P(y = 1|(w, c_i)) = \prod_{c_i \in c} \frac{1}{1 + e^{-e(w)e(c_i)}}$$

# Skip-Gram Method

- ▶ we treat words of a context independent from each other

$$P(y = 1|(c, w)) = \prod_{c_i \in c} P(y = 1|(w, c_i)) = \prod_{c_i \in c} \frac{1}{1 + e^{-e(w)e(c_i)}}$$

- ▶ loss in Skip-Gram is identical to that in CBoW

# Skip-Gram Method

- ▶ we treat words of a context independent from each other

$$P(y = 1|(c, w)) = \prod_{c_i \in c} P(y = 1|(w, c_i)) = \prod_{c_i \in c} \frac{1}{1 + e^{-e(w)e(c_i)}}$$

- ▶ loss in Skip-Gram is identical to that in CBoW
- ▶ we fine-tune parameters (word embeddings) using SGD

# CBoW and Skip-Gram

- ▶ CBoW loses the order information between context's elements



# CBoW and Skip-Gram

- ▶ CBoW loses the order information between context's elements
- ▶ CBoW allows the use of variable-length contexts

# CBoW and Skip-Gram

- ▶ CBoW loses the order information between context's elements
- ▶ CBoW allows the use of variable-length contexts
- ▶ Skip-Gram decouples the dependence between context elements even further than CBoW

# CBoW and Skip-Gram

- ▶ CBoW loses the order information between context's elements
- ▶ CBoW allows the use of variable-length contexts
- ▶ Skip-Gram decouples the dependence between context elements even further than CBoW
- ▶ Skip-Gram treats each context's element as an independent context

# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph

# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph
- ▶ window-based context: for a word we consider all words in a window of length  $2m + 1$  centred on the word as context

# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph
- ▶ window-based context: for a word we consider all words in a window of length  $2m + 1$  centred on the word as context
- ▶ the size of window: large windows tend to capture more topical similarities and small windows capture syntactic similarities

# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph
- ▶ window-based context: for a word we consider all words in a window of length  $2m + 1$  centred on the word as context
- ▶ the size of window: large windows tend to capture more topical similarities and small windows capture syntactic similarities
- ▶ contexts can be defined based on text units, for example, words that occur in the same sentence, paragraph, or text

# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph
- ▶ window-based context: for a word we consider all words in a window of length  $2m + 1$  centred on the word as context
- ▶ the size of window: large windows tend to capture more topical similarities and small windows capture syntactic similarities
- ▶ contexts can be defined based on text units, for example, words that occur in the same sentence, paragraph, or text
- ▶ contexts can also be defined based on syntax of a sentence for example using parse tree or dependency tree of a sentence



# Context Matters!

- ▶ context of a word translates to its surrounding words in a sentence or paragraph
- ▶ window-based context: for a word we consider all words in a window of length  $2m + 1$  centred on the word as context
- ▶ the size of window: large windows tend to capture more topical similarities and small windows capture syntactic similarities
- ▶ contexts can be defined based on text units, for example, words that occur in the same sentence, paragraph, or text
- ▶ contexts can also be defined based on syntax of a sentence for example using parse tree or dependency tree of a sentence
- ▶ contexts of a word can be foreign words that are aligned to the word in multilingual corpora

# Word2Vec Software Package

- ▶ the implementations of CBoW and Skip-Gram methods exist in Word2Vec
- ▶ <https://code.google.com/archive/p/word2vec/>

# Word2Vec Software Package

- ▶ the implementations of CBoW and Skip-Gram methods exist in Word2Vec
- ▶ <https://code.google.com/archive/p/word2vec/>
- ▶ Skip-Gram is more effective in practice

# GLoVe (Pennington et al., 2014)

- ▶ GloVe: Global Vectors for Word Representation
- ▶ GloVe is an unsupervised method for obtaining word embeddings
- ▶ GloVe aims at reconciling the advantages of corpus-wide co-occurrence counts and local context windows
- ▶ <https://nlp.stanford.edu/projects/glove/>

# Evaluating Word Embeddings

► intrinsic

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic
  - ▶ on a downstream NLP task, we compare the performance of two models that differ only in the word embeddings they use



# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic
  - ▶ on a downstream NLP task, we compare the performance of two models that differ only in the word embeddings they use
  - ▶ named entity recognition (NER): accuracy

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic
  - ▶ on a downstream NLP task, we compare the performance of two models that differ only in the word embeddings they use
  - ▶ named entity recognition (NER): accuracy
  - ▶ machine translation (MT): BLEU score

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic
  - ▶ on a downstream NLP task, we compare the performance of two models that differ only in the word embeddings they use
  - ▶ named entity recognition (NER): accuracy
  - ▶ machine translation (MT): BLEU score
  - ▶ summarization: ROUGE score

# Evaluating Word Embeddings

- ▶ intrinsic
  - ▶ word similarity tasks
  - ▶ word analogy tasks
- ▶ extrinsic
  - ▶ on a downstream NLP task, we compare the performance of two models that differ only in the word embeddings they use
  - ▶ named entity recognition (NER): accuracy
  - ▶ machine translation (MT): BLEU score
  - ▶ summarization: ROUGE score
  - ▶ information retrieval (IR): precision - recall - F1 score

# Word Similarity Tasks

- ▶ similar words should have similar representations
- ▶ dataset:  
<http://alfonseca.org/eng/research/wordsim353.html>

# Word Similarity Tasks

- ▶ similar words should have similar representations
- ▶ dataset:  
`http://alfonseca.org/eng/research/wordsim353.html`
- ▶ word-1 and word-2  $\rightarrow$  similarity score  $\in [0, 10]$

# Word Similarity Tasks

- ▶ similar words should have similar representations
- ▶ dataset:  
<http://alfonseca.org/eng/research/wordsim353.html>
- ▶ word-1 and word-2  $\rightarrow$  similarity score  $\in [0, 10]$
- ▶ our function  $f()$  should map two words to a similarity score using the distance between the word vectors of the words

# Word Similarity Tasks

- ▶ similar words should have similar representations
- ▶ dataset:  
<http://alfonseca.org/eng/research/wordsim353.html>
- ▶ word-1 and word-2  $\rightarrow$  similarity score  $\in [0, 10]$
- ▶ our function  $f()$  should map two words to a similarity score using the distance between the word vectors of the words
- ▶ cosine similarity

$$\text{sim}(w_i, w_j) = \frac{e(w_i)e(w_j)}{\|e(w_i)\|^2\|e(w_j)\|^2}$$



# Word Analogy Tasks

- ▶ A is to B as C to ?

# Word Analogy Tasks

- ▶ A is to B as C to ?
- ▶ Germany is to Berlin as France is to ?
- ▶ dataset: `http://download.tensorflow.org/data/questions-words.txt`

# Word Analogy Tasks

- ▶ A is to B as C to ?
- ▶ Germany is to Berlin as France is to ?
- ▶ dataset: <http://download.tensorflow.org/data/questions-words.txt>
- ▶ capital-common-countries
  - ▶ Athens Greece Baghdad → Iraq
  - ▶ Athens Greece Berlin → Germany

# Word Analogy Tasks

- ▶ A is to B as C to ?
- ▶ Germany is to Berlin as France is to ?
- ▶ dataset: <http://download.tensorflow.org/data/questions-words.txt>
- ▶ capital-common-countries
  - ▶ Athens Greece Baghdad → Iraq
  - ▶ Athens Greece Berlin → Germany
- ▶ family
  - ▶ boy girl brother → sister
  - ▶ brother sister dad → mom

# Word Analogy Tasks

- ▶ A is to B as C to ?
- ▶ Germany is to Berlin as France is to ?
- ▶ dataset: <http://download.tensorflow.org/data/questions-words.txt>
- ▶ capital-common-countries
  - ▶ Athens Greece Baghdad → Iraq
  - ▶ Athens Greece Berlin → Germany
- ▶ family
  - ▶ boy girl brother → sister
  - ▶ brother sister dad → mom
- ▶ currency, adj-to-adverb, comparative, ...

# Finding the Prototype of a Group of Words

- ▶ if we have a group of words  $g = \{w_1, w_2, \dots, w_k\}$

# Finding the Prototype of a Group of Words

- ▶ if we have a group of words  $g = \{w_1, w_2, \dots, w_k\}$
- ▶ the prototype of this group can be computed as follows

$$\text{proto}(g) = \frac{1}{k} \sum_{i \in 1..k} e(w_i)$$

where  $e$  maps a word to its embeddings

# Finding Similar Words

- ▶ words found by embedding-based similarities can be filtered with other types of word similarities



# Finding Similar Words

- ▶ words found by embedding-based similarities can be filtered with other types of word similarities

# Finding Similar Words

- ▶ words found by embedding-based similarities can be filtered with other types of word similarities

# Finding Similar Words

- ▶ cosine similarity

$$\text{sim}(w_i, w_j) = \frac{e(w_i)e(w_j)}{\|e(w_i)\|^2\|e(w_j)\|^2}$$

- ▶ words found by embedding-based similarities can be filtered with other types of word similarities

# Short Text Similarities

►  $d_1 = \{w_1^1, w_2^1, \dots, w_m^1\}$  and  $d_2 = \{w_1^2, w_2^2, \dots, w_n^2\}$

# Short Text Similarities

►  $d_1 = \{w_1^1, w_2^1, \dots, w_m^1\}$  and  $d_2 = \{w_1^2, w_2^2, \dots, w_n^2\}$



$$\text{sim}(d_1, d_2) = \frac{1}{m \cdot n} \sum_{i=1}^m \sum_{j=1}^n \cos(e(w_i^1), e(w_j^2))$$

# Pre-Trained Embeddings

- ▶ Word2Vec
  - ▶ trained on Google News (100 billion tokens)
- ▶ GloVe
  - ▶ trained on Wikipedia (6 billion tokens)
  - ▶ trained on CommonCrawl (42 and 840 billion tokens)
  - ▶ trained on Twitter (27 million tokens)
- ▶ many other pre-trained embeddings for different languages
  - ▶ <https://fasttext.cc/docs/en/crawl-vectors.html>

# Practical Hints

- ▶ always try out different word embeddings (consider them as a hyperparameter)

# Practical Hints

- ▶ always try out different word embeddings (consider them as a hyperparameter)
- ▶ results may vary drastically with different embeddings



# Practical Hints

- ▶ always try out different word embeddings (consider them as a hyperparameter)
- ▶ results may vary drastically with different embeddings
- ▶ consider source of corpora used to train word embeddings (larger is not always better, a smaller but more domain-focused can be more effective)

# Practical Hints

- ▶ always try out different word embeddings (consider them as a hyperparameter)
- ▶ results may vary drastically with different embeddings
- ▶ consider source of corpora used to train word embeddings (larger is not always better, a smaller but more domain-focused can be more effective)
- ▶ consider what contexts were used to define similarities

# Practical Hints

- ▶ always try out different word embeddings (consider them as a hyperparameter)
- ▶ results may vary drastically with different embeddings
- ▶ consider source of corpora used to train word embeddings (larger is not always better, a smaller but more domain-focused can be more effective)
- ▶ consider what contexts were used to define similarities
- ▶ it's better to use the same tokenization and text normalization methods that were used for creating word embeddings

# Embedding Layers in PyTorch

- ▶ an embedding layer (a.k.a lookup table) maps a sequence of word IDs to a sequence of embedding vectors

# Embedding Layers in PyTorch

- ▶ an embedding layer (a.k.a lookup table) maps a sequence of word IDs to a sequence of embedding vectors
- ▶ it uses a lookup table, which is a matrix with size  $|V| \times d_{emb}$

# Embedding Layers in PyTorch

- ▶ an embedding layer (a.k.a lookup table) maps a sequence of word IDs to a sequence of embedding vectors
- ▶ it uses a lookup table, which is a matrix with size  $|V| \times d_{emb}$
- ▶ the  $i$ 'th row of this matrix contains embeddings of  $i$ 'th word in vocabulary  $V$

# Embedding Layers in PyTorch

- ▶ an embedding layer (a.k.a lookup table) maps a sequence of word IDs to a sequence of embedding vectors
- ▶ it uses a lookup table, which is a matrix with size  $|V| \times d_{emb}$
- ▶ the  $i$ 'th row of this matrix contains embeddings of  $i$ 'th word in vocabulary  $V$
- ▶ so the only thing we need to do is to replace a word with its index in vocabulary  $\rightarrow$  a dictionary does this easily (`word_to_ix`)

# Embedding Layers in PyTorch

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim

torch.manual_seed(1)

word_to_ix = {"hello": 0, "world": 1}

embeds = nn.Embedding(2, 5)

lookup_tensor = torch.tensor([word_to_ix["hello"]], dtype=torch.long)

hello_embed = embeds(lookup_tensor)

print(hello_embed)
```



# Embedding Layers in PyTorch

```
import torch
import torch.nn as nn

class EmbeddingLayer(nn.Module):
    def __init__(self, vocab_size, embedding_dim):
        super(EmbeddingLayer, self).__init__()
        self.embeddings = nn.Embedding(vocab_size, embedding_dim)

    def forward(self, inputs):
        embeds = self.embeddings(inputs)
        return embeds
```

# Embedding Layers in PyTorch

```
import torch
import torch.nn as nn

class EmbeddingLayer(nn.Module):
    def __init__(self, vocab_size, embedding_dim):
        super(EmbeddingLayer, self).__init__()
        self.embeddings = nn.Embedding(vocab_size, embedding_dim)

    def forward(self, inputs):
        embeds = self.embeddings(inputs)
        return embeds

if __name__ == '__main__':
    model = EmbeddingLayer(vocab_size=10, emb_size=3)
    inputs = [1, 5]
    inputs_tenseor = torch.tensor(inputs, dtype=torch.long)

    emb_vectors = model(inputs_tenseor)

    print(f"the_shape_of_emb_vectors_is_{emb_vectors.shape}")
```

# Using Pretrained Embedding Layers in PyTorch

```
import torch
import torch.nn as nn

class EmbeddingLayer(nn.Module):
    def __init__(self, vocab_size, embedding_dim, weights_matrix):
        super(EmbeddingLayer, self).__init__()
        self.embeddings = nn.Embedding(vocab_size, embedding_dim)

        self.embeddings.load_state_dict({'weight': weights_matrix})

    def forward(self, inputs):
        embeds = self.embeddings(inputs)
        return embeds
```

# Freezing Embedding Layers in PyTorch

```
import torch
import torch.nn as nn

class EmbeddingLayer(nn.Module):
    def __init__(self, vocab_size, embedding_dim, weights_matrix, freeze=False):
        super(EmbeddingLayer, self).__init__()
        self.embeddings = nn.Embedding(vocab_size, embedding_dim)

        self.embeddings.load_state_dict({'weight': weights_matrix})

        if freeze:
            self.embeddings.weight.requires_grad = False

    def forward(self, inputs):
        embeds = self.embeddings(inputs)
        return embeds
```

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets
  - ▶ “cat” is more similar to “tiger” than to “dog” as they both as felines

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets
  - ▶ “cat” is more similar to “tiger” than to “dog” as they both as felines
- ▶ many of the trivial properties of words are ignored because people are less likely to mention known information than they are to mention novel one



# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets
  - ▶ “cat” is more similar to “tiger” than to “dog” as they both as felines
- ▶ many of the trivial properties of words are ignored because people are less likely to mention known information than they are to mention novel one
  - ▶ when people talk about “white sheep”, they will likely prefer to use only “sheep”

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets
  - ▶ “cat” is more similar to “tiger” than to “dog” as they both as felines
- ▶ many of the trivial properties of words are ignored because people are less likely to mention known information than they are to mention novel one
  - ▶ when people talk about “white sheep”, they will likely prefer to use only “sheep”
  - ▶ for “black sheep”, they are very likely to use color information “black sheep”

# Limitations

- ▶ the algorithms discussed provide very little control over the kind of similarity they include
  - ▶ “cat” is more similar to “dog” than to “tiger” as they both are pets
  - ▶ “cat” is more similar to “tiger” than to “dog” as they both as felines
- ▶ many of the trivial properties of words are ignored because people are less likely to mention known information than they are to mention novel one
  - ▶ when people talk about “white sheep”, they will likely prefer to use only “sheep”
  - ▶ for “black sheep”, they are very likely to use color information “black sheep”
  - ▶ a model trained on texts only might easily misled by this

# Limitations

- ▶ text corpora can easily be biased for better or worse → so word embeddings can become biased too

# Limitations

- ▶ text corpora can easily be biased for better or worse → so word embeddings can become biased too
  - ▶ gender and racial biases are very common

# Limitations

- ▶ text corpora can easily be biased for better or worse → so word embeddings can become biased too
  - ▶ gender and racial biases are very common
- ▶ these word vectors are context independent

# Limitations

- ▶ text corpora can easily be biased for better or worse → so word embeddings can become biased too
  - ▶ gender and racial biases are very common
- ▶ these word vectors are context independent
  - ▶ in reality, there is no such a thing to have context independent word meaning

# Limitations

- ▶ text corpora can easily be biased for better or worse → so word embeddings can become biased too
  - ▶ gender and racial biases are very common
- ▶ these word vectors are context independent
  - ▶ in reality, there is no such a thing to have context independent word meaning
  - ▶ some words have multiple sense e.g., “bank” may refer to a financial institution or to the side of a river



- ▶ common features used for converting textual data into numerical vectors
- ▶ basics of word embeddings
  - ▶ how to get them?
  - ▶ where to use them?
  - ▶ how to use them?
- ▶ limitations

Thank You!