

Introduction

Dense Text Representations

Nils Reimers

Author of Sentence Transformers

www.SBERT.net

Content

- **Part 1: Introduction**

- Background
- Applications

- Part 2: Basic Training Methods

- Part 3: Advanced Training Methods

Why Dense Representations?

- Text = Sentence, Paragraph, Document
- Traditionally: Sparse, Lexical Representation (Bag-of-words)
 - Each word has its own dimension

How are you?
[0, 0, 1, 0, 0, 0, 1, 0, 1, ...]
 ↑ ↑ ↑
 How are you

- Issues:
 - Lexical gap: US, USA, United States
 - Ambiguous words
 - Same distance between all words
 - Word order not preserved

Dense Representation

- Formally: $f(\text{Text}) \rightarrow \mathbb{R}^n$
- n dimensional representation (embedding)
- Find function f such that semantically similar text is close

What does semantically similar mean?

- Find function f such that **semantically similar** text is close
- There cannot be ***universal** text representations*

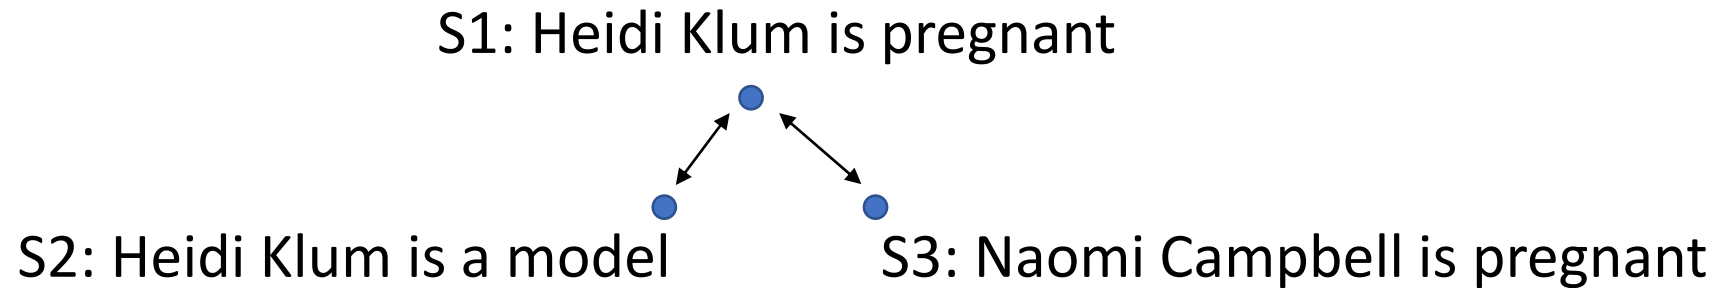
Nuclear energy is safe! ● \longleftrightarrow ● Nuclear energy is dangerous!



Semantically similar depends on the task!

What does semantically similar mean?

- Find function f such that **semantically similar** text is close
- There cannot be ***universal text representations***

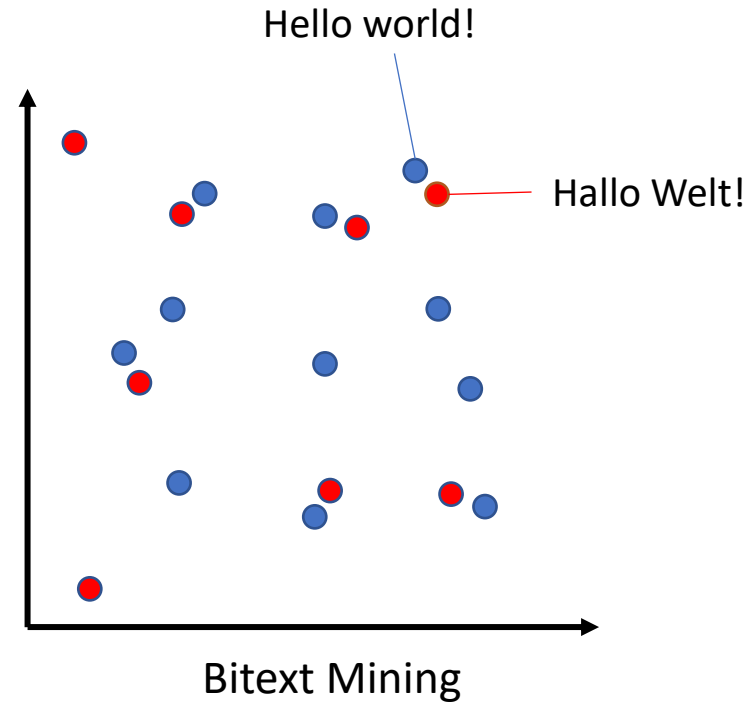


Semantically similar depends on the task!


Application - Clustering



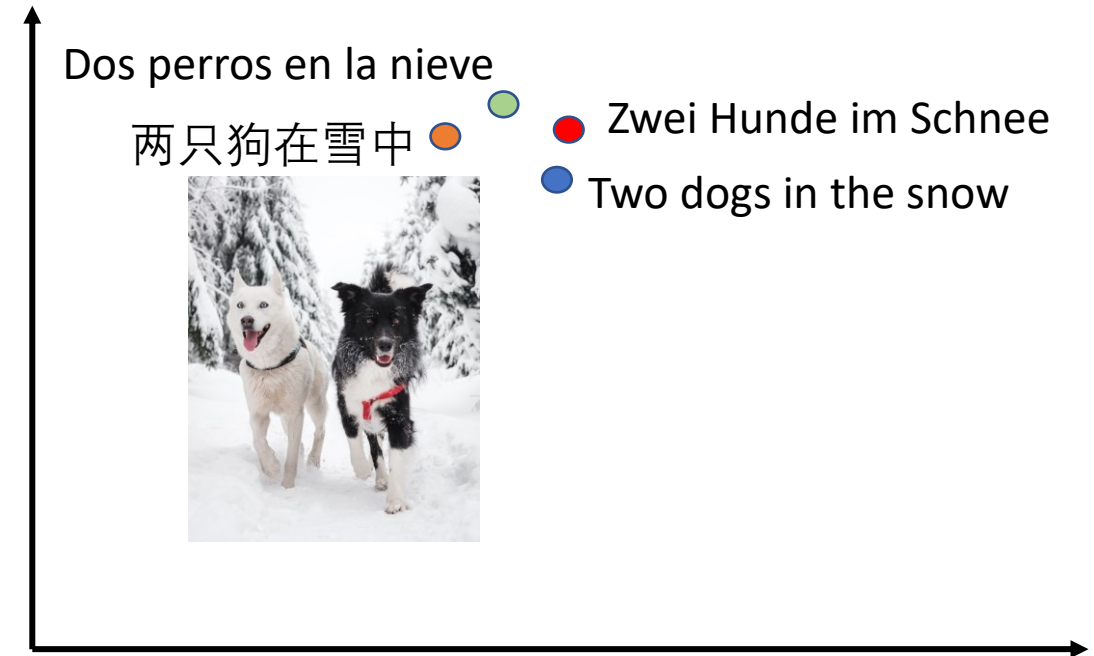
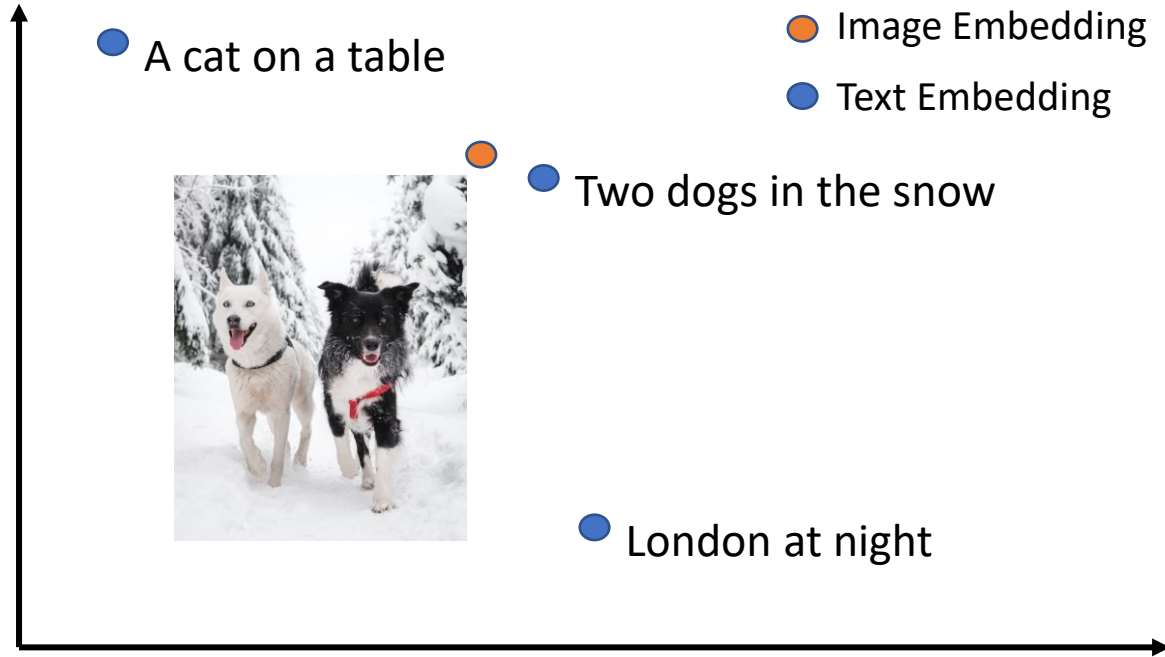
Application – Bitext Mining



Application – Search



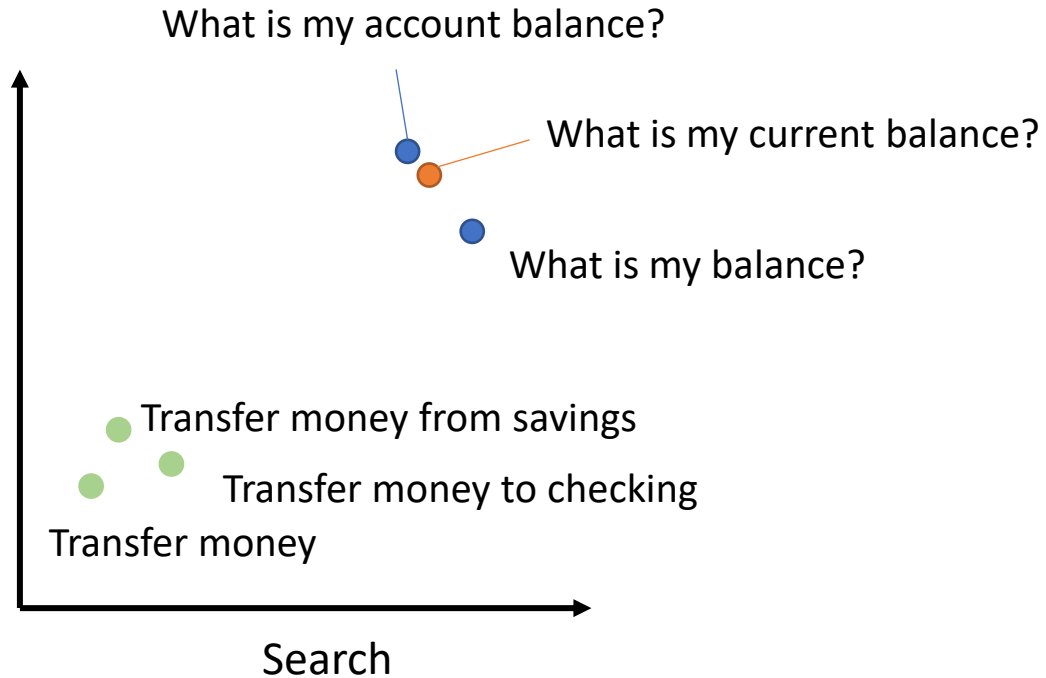
Multi-Modal Search



Zero-Shot Image Classification




Few-Shot Intent Classification




- Have some examples for every intent (checking balance, transfer money)
- New utterance => find closest example => use intent

Application – Automate E-Mail Support



- Find most common e-mails in your inbox
- Create canned responses for top 100 questions
- Train model on (email, response)



- For new email: What is the closest (canned) response?

Conclusion

- Map text/images/... to low dimensional dense vector space
- Semantically similar text should be close
 - No clear definition what “semantically similar” means => depends on task
- Many applications. Most promising:
 - Search
 - Mining – find related items

Basic Training Dense Text Representations

Nils Reimers

Author of Sentence Transformers

www.SBERT.net

Content

- Part 1: Introduction
- **Part 2: Basic Training Methods**
 - Basic training method
 - Loss functions
 - Improving training quality with hard negatives
- Part 3: Advanced Training Methods

Average Word Embeddings

- Simple baseline: Average word embeddings in a sentence

$$\begin{array}{c} W_1 \\ \begin{bmatrix} W_{11} \\ W_{12} \\ \vdots \\ W_{1n} \end{bmatrix} \end{array} + \begin{array}{c} W_2 \\ \begin{bmatrix} W_{21} \\ W_{22} \\ \vdots \\ W_{2n} \end{bmatrix} \end{array} + \dots + \begin{array}{c} W_n \\ \begin{bmatrix} W_{n1} \\ W_{n2} \\ \vdots \\ W_{nn} \end{bmatrix} \end{array} = \begin{array}{c} D \\ \begin{bmatrix} \frac{W_{11} + W_{21} + \dots + W_{n1}}{n} \\ \frac{W_{12} + W_{22} + \dots + W_{n2}}{n} \\ \vdots \\ \frac{W_{1n} + W_{2n} + \dots + W_{nn}}{n} \end{bmatrix} \end{array}$$

- Improvement: Weight word embeddings by tf-idf
 - Content words contribute more than stop words
- Fast and simple method
- Can be used with custom trained word embeddings

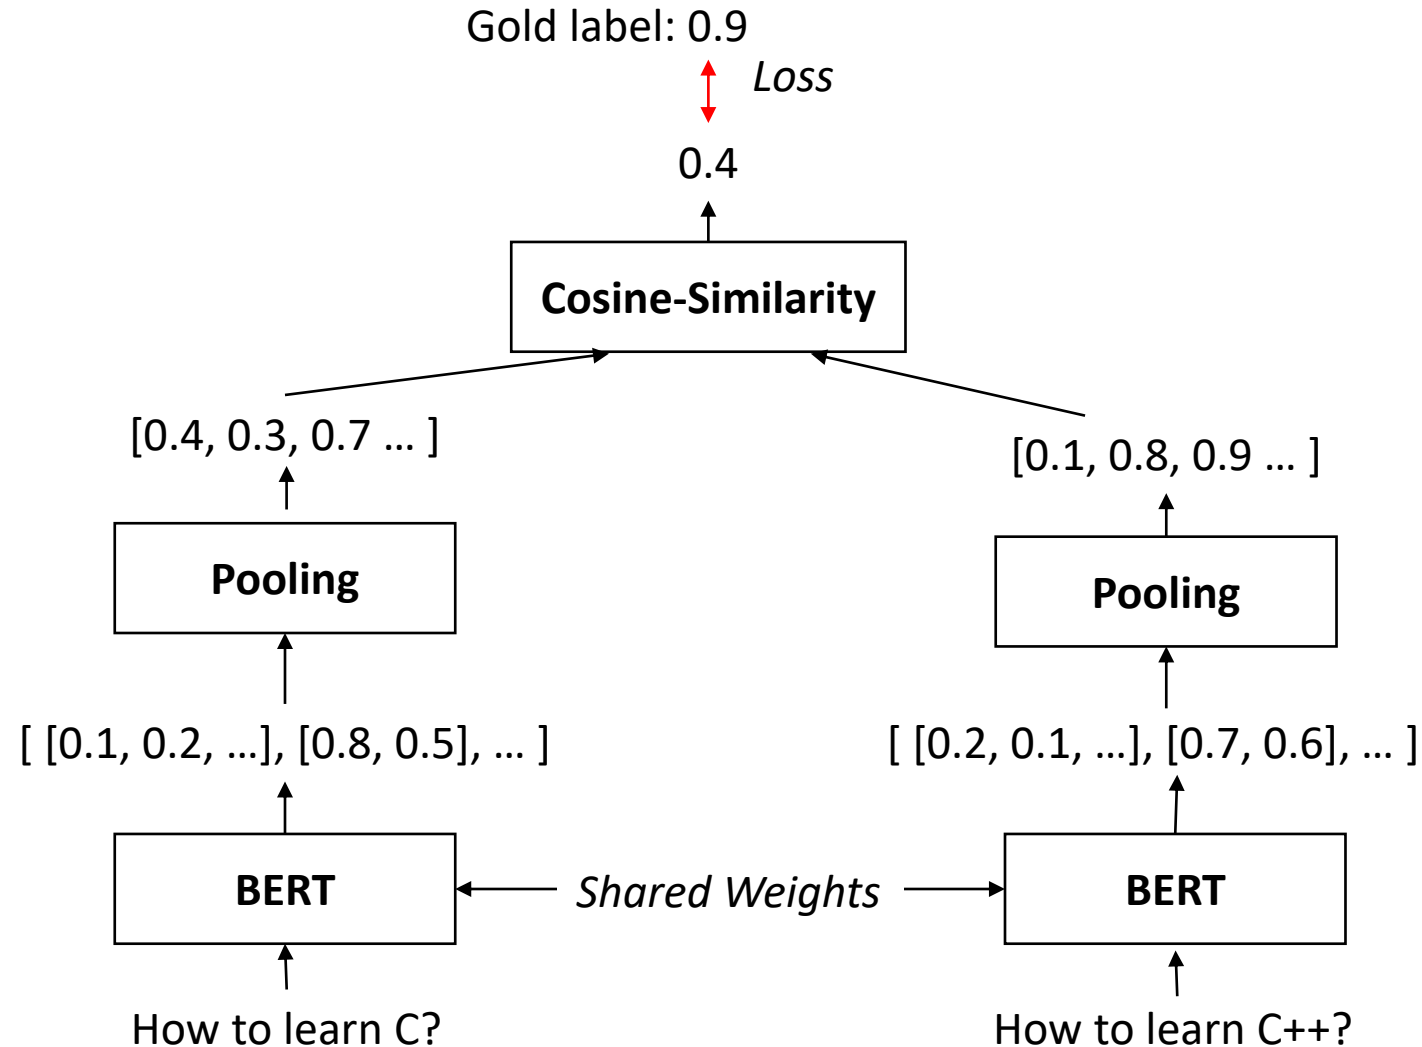
Can we just average BERT?

=> No!

Task	Avg. performance on 14 datasets
bert-base-uncased	48.5
Avg. word embeddings	51.1
Best unsupervised method based on BERT-large	60.8
SBERT (NLI data)	62.5
SBERT (large train data)	68.0

- BERT out-of-the-box performs badly
- Labeled / structured data important

Sentence BERT – Cosine Similarity Loss

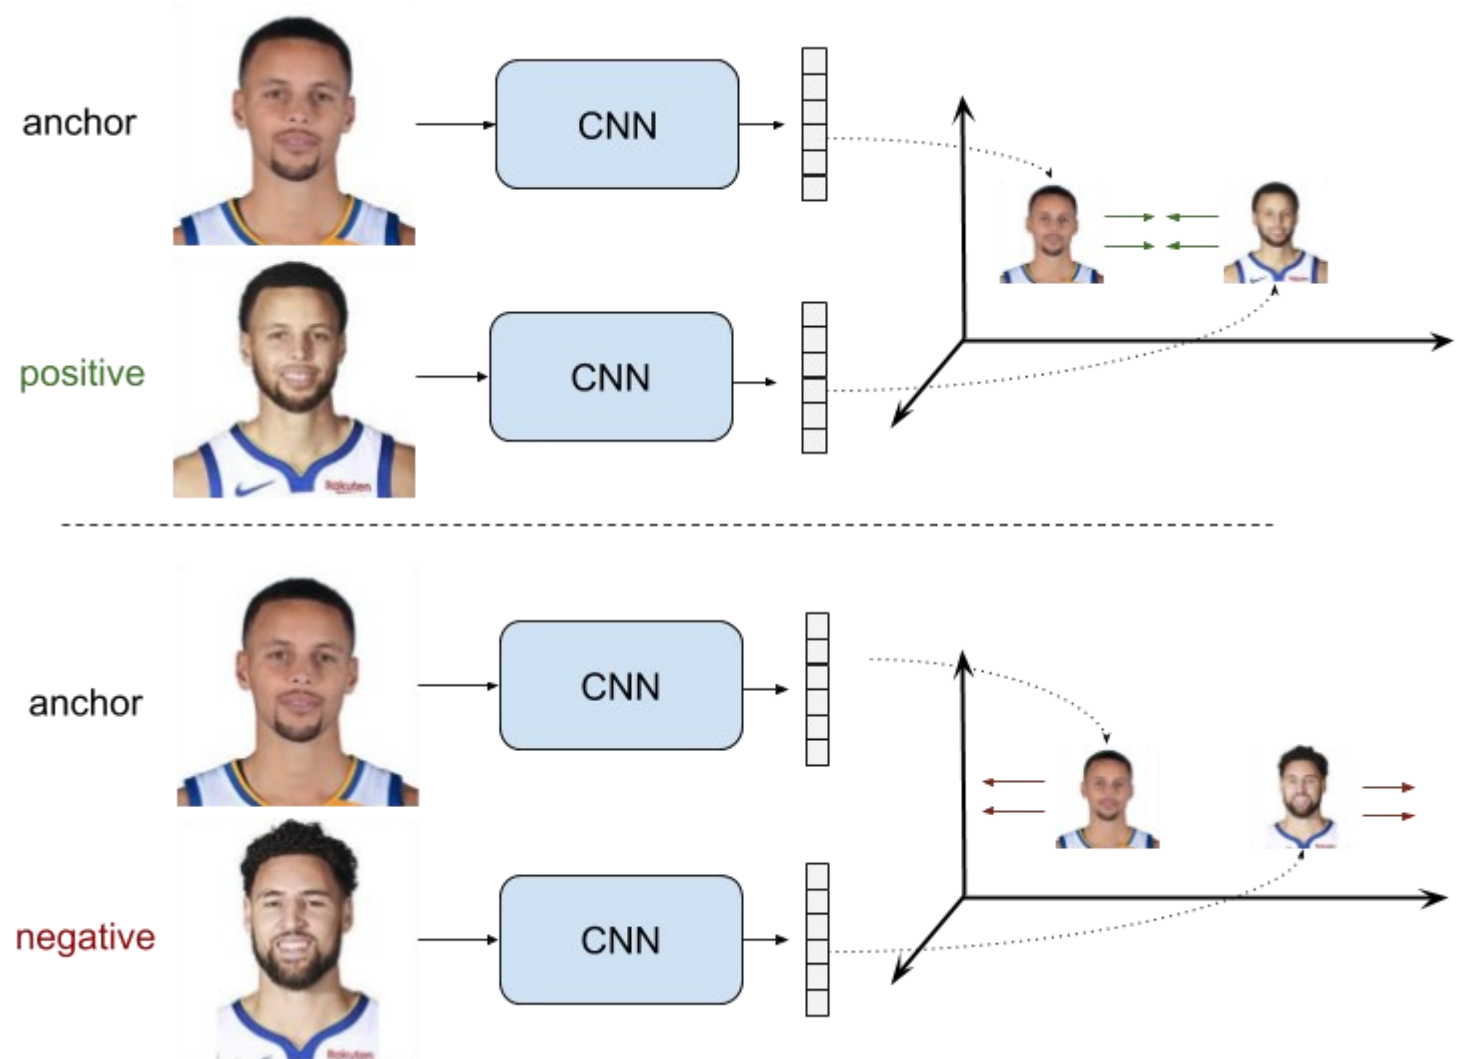


Loss Functions

- Various loss functions available to train vector space models:
 - CosineSimilarityLoss
 - SoftmaxLoss
 - Constrative- & OnlineConstrativeLoss
 - TripletLoss
 - Batch-[All | Hard | SemiHard | HardSoftMargin]-TripletLoss
 - MegaBatchMarginLoss
 - MultipleNegativesRankingLoss
 -


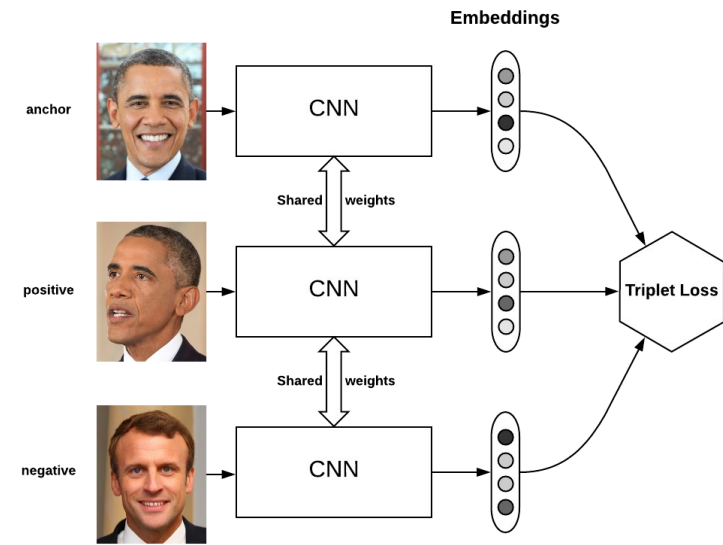

Contrastive Loss

- Positive pairs
 - Pull together
- Negative pairs
 - Push away




Triplet Loss

- Triplet loss
 - Requires (anchor, positive, negative)
 - Anchor & negative should be further away than anchor & positive
 - To work well: requires good triplets




Global And Local Structure of Vector Space

- Global structure: Relation of two random sentences
- Local structure: Relation of two similar sentences
- Loss function must optimize local and global structure
- Constrative / Triplet loss might only optimize the local structure



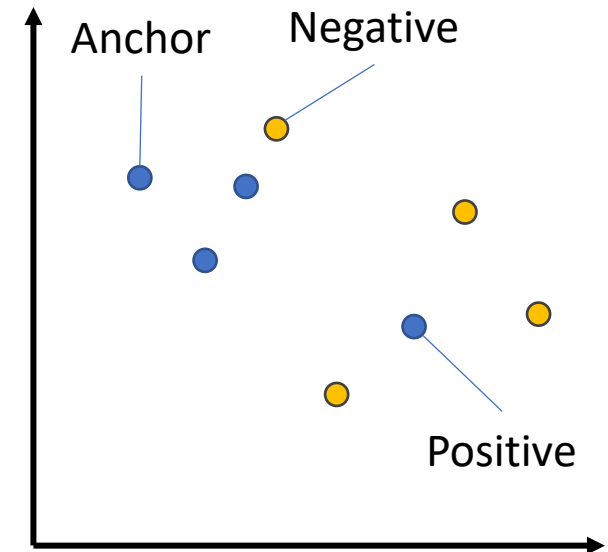
Bad global structure



Good global structure

Batch Hard Triplet Loss

- Have a batch with multiple examples for each class
- Take an element as anchor
 - Take the farthest away element as positive
 - Take the closest element as negative
 - Train with triplet loss
- Loss automatically creates hard triplets



Multiple Negative Ranking Loss

- Have positive pairs:


(a_1, p_1)

(a_2, p_2)

(a_3, p_3)

- Examples:

- (query, answer-passage)
- (question, duplicate_question)
- (paper title, cited paper title)




- (a_i, p_i) should be close in vector space and (a_i, p_j) should be distant in vector space ($i \neq j$)
 - Unlikely that e.g. two randomly selected questions are similar
- Computed as ranking loss with Cross-Entropy:
 - Given a_1 , which is the right answer out of $[p_1, p_2, p_3]$?
 - Compute scores: $[s(a_1, p_1), s(a_1, p_2), s(a_1, p_3)]$
 - Cross-Entropy loss with gold label: $[1, 0, 0]$
- Also called “training with in-batch negatives”, InfoNCE or NTXentLoss

Multiple Negative Ranking Loss

Hard Negatives

- Larger batch size => task more difficult => better results
 - Given query, which of the 10 passages provide the answer?
 - Given query, which of the 1k passages provide the answer?



Multiple Negative Ranking Loss


Hard Negatives

- Train with tuples:

(a_1, p_1, n_1)

(a_2, p_2, n_2)


- n_i should be similar to p_i but not match with a_i
- Bad example:
 - a: How many people live in London?
 - p: Around 9 million people live in London
 - n: London has a population of 9 million people.
- Good example:
 - a: How many people live in London?
 - p: Around 9 million people live in London
 - n: Around 1 million people live in Birmingham, second to London.



How to find hard-negatives?

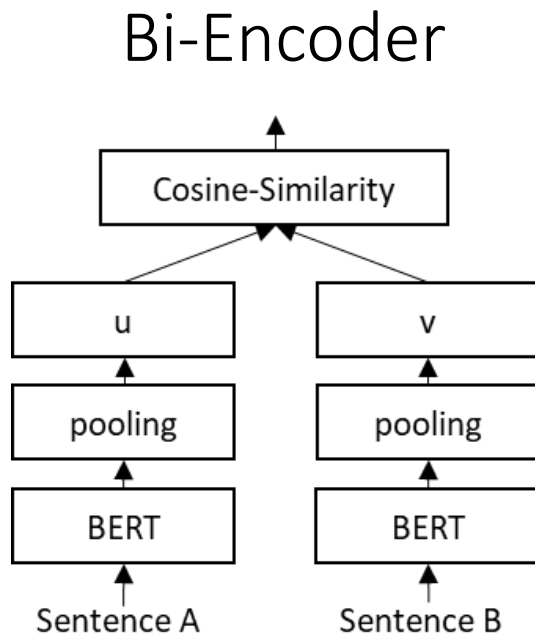
- Quality of hard-negatives significantly improves the performance
- Finding good hard negatives not easy
- Strategy 1: Exploit structure in your data
 - Citation graph: (Title, Cited_Paper, Paper_Cited_by_Cited_Paper)
 - Q&A: (Question, Answer with most stars, Answer with less stars)
- Strategy 2: Mine hard negatives (simple):
 - Use BM25 to find top-100 most similar texts to anchor / positive
 - Select one of these randomly

Issue with BM25-Negatives




- BM25 negatives not necessarily hard negatives for the Dense Representations (DR)


Bi- vs Cross-Encoders




Cross-Encoder



Cross vs. Bi-Encoders in STSb (English)



How to find hard-negatives?



How big is the improvement?

Approach	MRR@10 on MSMARCO Dev
Random negatives	26.1
BM25 negatives	29.9
Mined hard negatives w/o denoising	26.0
Mined hard negatives with denoising	36.4

Source: <https://arxiv.org/pdf/2007.00808.pdf> and <https://arxiv.org/pdf/2010.08191.pdf>

Conclusion

- Many different loss functions available
 - In many cases, MultipleNegativeRankingLoss work well
- Adding hard negatives improves performance for search
 - but not for clustering!
- Finding hard negatives not trivial
 - Usage of powerful cross-encoders to mine those

Advanced Training Dense Text Representations

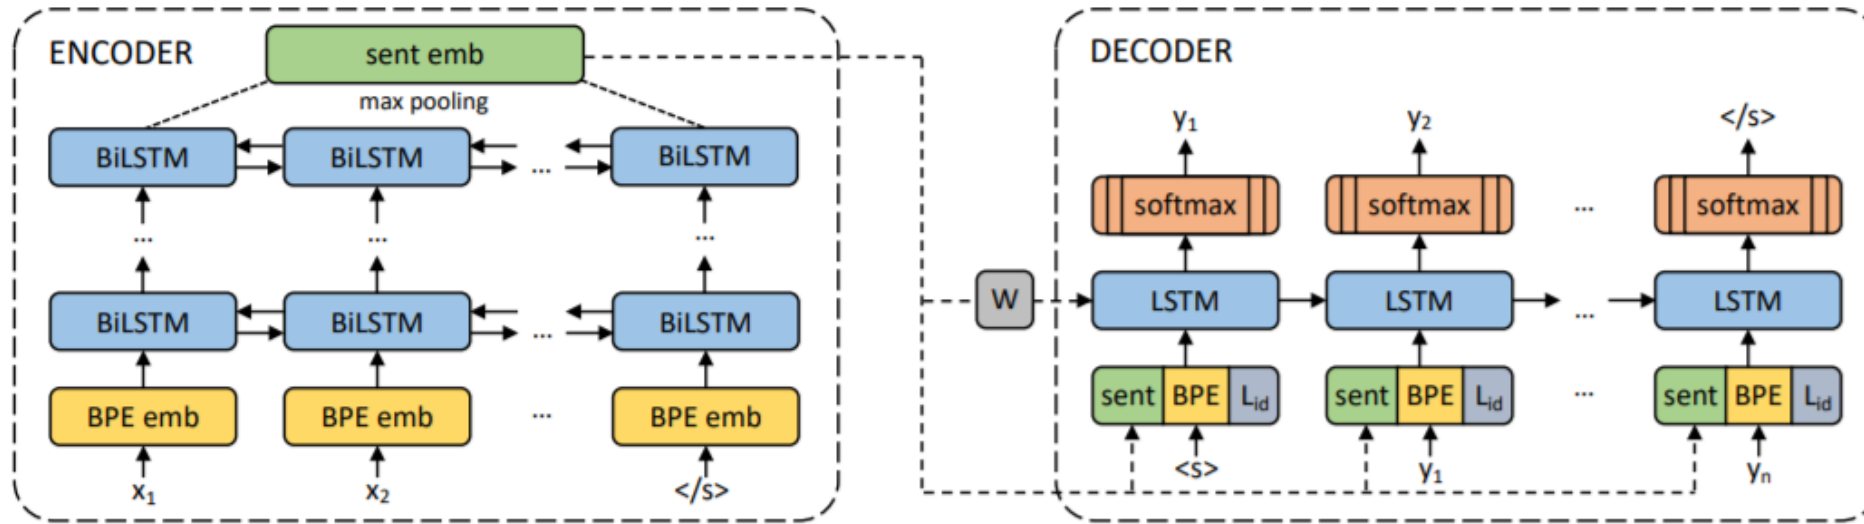
Nils Reimers

www.SBERT.net

Content

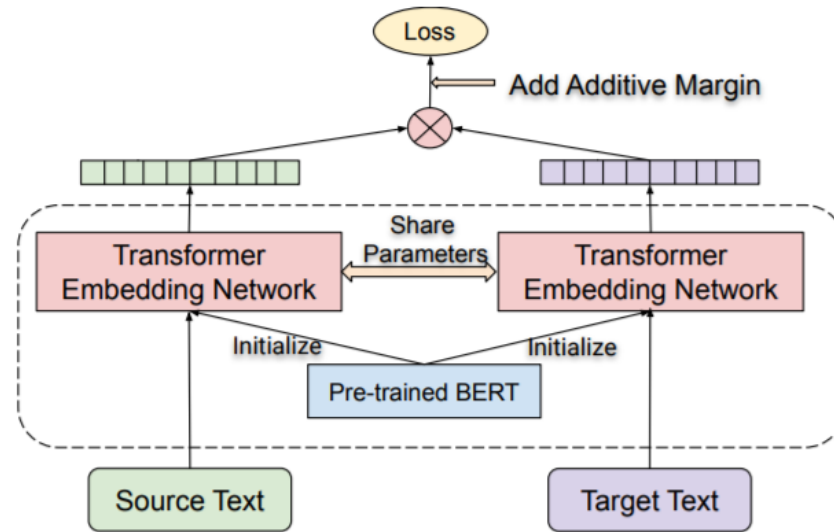
- Part 1: Introduction
- Part 2: Basic Training Methods
- **Part 3: Advanced Training Methods**
 - Multilingual Text Embeddings
 - Data Augmentation using Cross-Encoders
 - Unsupervised Embedding Learning
 - Neural Search / BEIR

Multilingual Sentence Embeddings: LASER



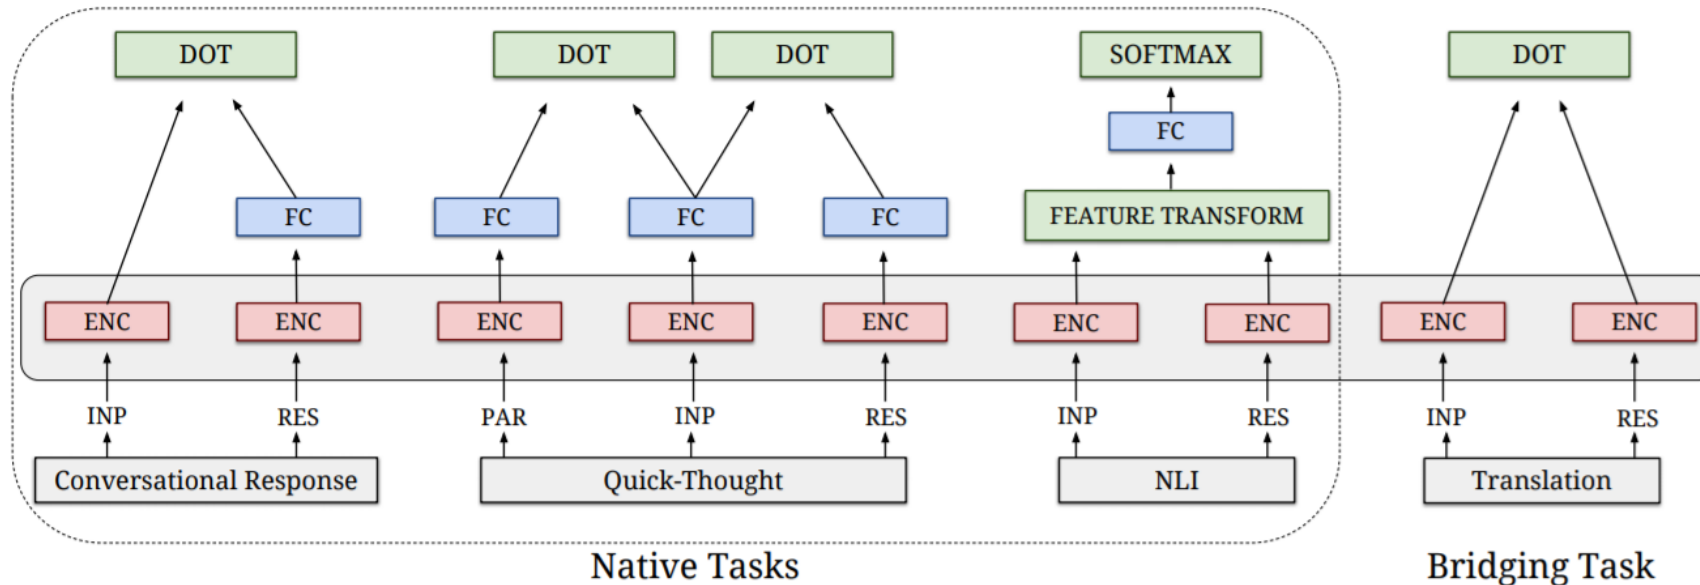
- Use output of encoder from translation system
- Issues:
 - Cannot control what type of embeddings are learned
 - Works poorly on identifying similar sentences

Multilingual Sentence Embeddings: LaBSE



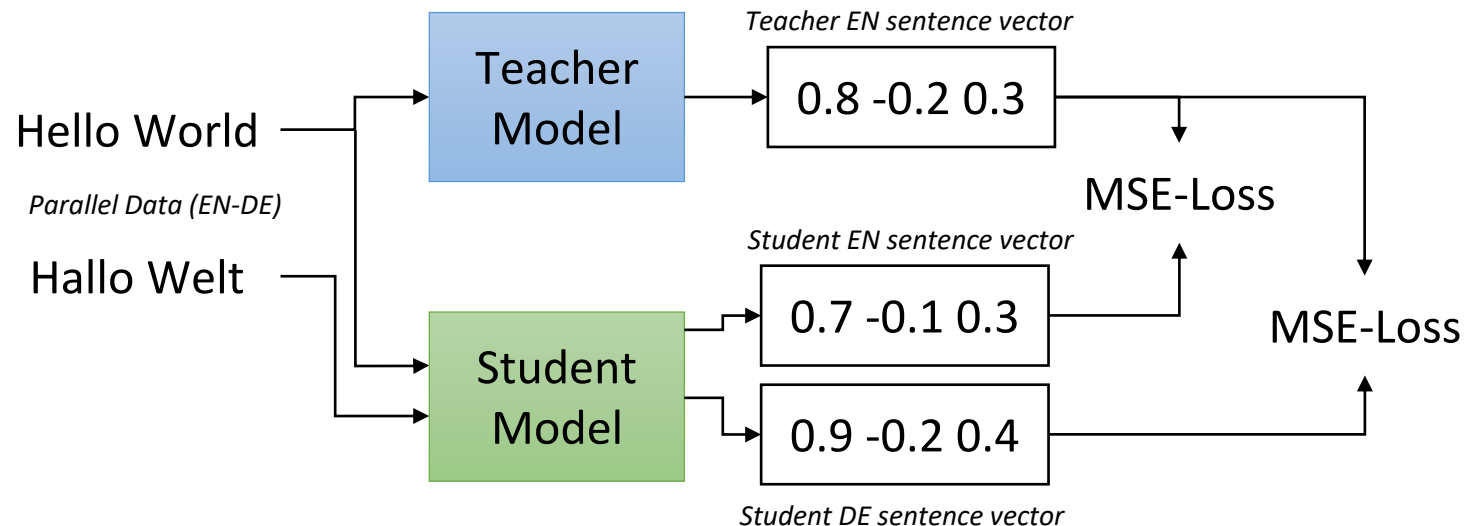
- Translation ranking task
- Issues:
 - Cannot control what type of embeddings are learned
 - Works poorly on identifying similar sentences

Multilingual Sentence Embeddings: mUSE



- Multi-task setup with bridging task
- Issues:
 - Getting bridging task right is challenging + requires large batch sizes
 - Hard to extend model afterwards to new languages

Multilingual Knowledge Distillation



- Given:
 - Teacher sentence embedding model T (e.g. SBERT trained on English STS)
 - Parallel sentence data $((s_1, t_1), \dots, (s_n, t_n))$
 - Student model S with multilingual vocabulary (e.g. XLM-R + Mean Pooling)
- Train student S such that:

$$S(s_i) \approx T(s_i)$$

$$S(t_i) \approx T(s_i)$$

Results – Multilingual Semantic Similarity

Model	EN-AR	EN-DE	EN-TR	EN-ES	EN-FR	EN-IT	EN-NL	Avg.
mBERT mean	16.7	33.9	16.0	21.5	33.0	34.0	35.6	27.2
XLM-R mean	17.4	21.3	9.2	10.9	16.6	22.9	26.0	17.8
mBERT-nli-stsb	30.9	62.2	23.9	45.4	57.8	54.3	54.1	46.9
XLM-R-nli-stsb	44.0	59.5	42.4	54.7	63.4	59.4	66.0	55.6
Knowledge Distillation								
mBERT \leftarrow SBERT-nli-stsb	77.2	78.9	73.2	79.2	78.8	78.9	77.3	77.6
DistilmBERT \leftarrow SBERT-nli-stsb	76.1	77.7	71.8	77.6	77.4	76.5	74.7	76.0
XLM-R \leftarrow SBERT-nli-stsb	77.8	78.9	74.0	79.7	78.5	78.9	77.7	77.9
XLM-R \leftarrow SBERT-paraphrases	82.3	84.0	80.9	83.1	84.9	86.3	84.5	83.7
Other Systems								
LASER	66.5	64.2	72.0	57.9	69.1	70.8	68.5	67.0
mUSE	79.3	82.1	75.5	79.6	82.6	84.5	84.1	81.1
LaBSE	74.5	73.8	72.0	65.5	77.0	76.9	75.1	73.5

- Training on English-only insufficient
- LASER & LaBSE perform badly

Bitext Mining

- Given two corpora: Find parallel (translated) sentences

Model	DE-EN	FR-EN	RU-EN	ZH-EN	Avg.
mBERT mean	44.1	47.2	38.0	37.4	41.7
XLM-R mean	5.2	6.6	22.1	12.4	11.6
mBERT-nli-stsb	38.9	39.5	26.4	30.2	33.7
XLM-R-nli-stsb	44.0	51.0	51.5	44.0	47.6
Knowledge Distillation					
XLM-R \leftarrow SBERT-nli-stsb	86.8	84.4	86.3	85.1	85.7
XLM-R \leftarrow SBERT-paraphrase	90.8	87.1	88.6	87.8	88.6
Other systems					
mUSE	88.5	86.3	89.1	86.9	87.7
LASER	95.4	92.4	92.3	91.7	93.0
LaBSE	95.9	92.5	92.4	93.0	93.5

Table 3: F_1 score on the BUCC bitext mining task.

- LASER & LaBSE better than mUSE & Knowledge Distillation
- Issue with mUSE & KD: They find similar sentences, that are not perfect translations

Data Efficiency

Dataset	#DE	EN-DE	#AR	EN-AR
XLM-R mean	-	21.3	-	17.4
XLM-R-nli-stsb	-	59.5	-	44.0
MUSE Dict	101k	75.8	27k	68.8
Wikitles Dict	545k	71.4	748k	67.9
MUSE + Wikitles	646k	76.0	775k	69.1

Dataset size	EN-DE	EN-AR
XLM-R mean	21.3	17.4
XLM-R-nli-stsb	59.5	44.0
1k	71.5	48.4
5k	74.5	59.6
10k	77.0	69.5
25k	80.0	70.2
Full TED2020	80.4	78.0


Table 6: Performance on STS 2017 dataset when trained with reduced TED2020 dataset sizes.

Knowledge Distillation vs. Training on Target Language

Model	KO-KO
LASER	68.44
mUSE	76.32
Trained on KorNLI & KorSTS	
Korean RoBERTa-base	80.29
Korean RoBERTa-large	80.49
XLM-R	79.19
XLM-R-large	81.84
Multiling. Knowledge Distillation	
XLM-R \leftarrow SBERT-nli-stsb	81.47
XLM-R-large \leftarrow SBERT-large-nli-stsb	83.00

Table 7: Spearman rank correlation on Korean STS-benchmark test-set (Ham et al., 2020).

Language Bias



Model	Expected Score	Actual Score	Difference
LASER	69.5	68.6	-0.92
mUSE	81.7	81.6	-0.19
LaBSE	74.4	73.1	-1.29
XLM-R \leftarrow SBERT-paraphrases	84.0	83.9	-0.11



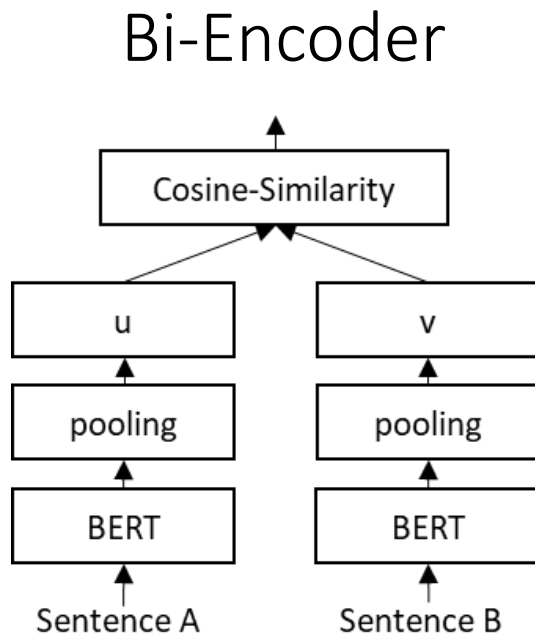
- Preference of certain language combinations
- Language bias impacts performance negatively on multilingual pools
- LASER and LaBSE with strong language bias

Augmented SBERT

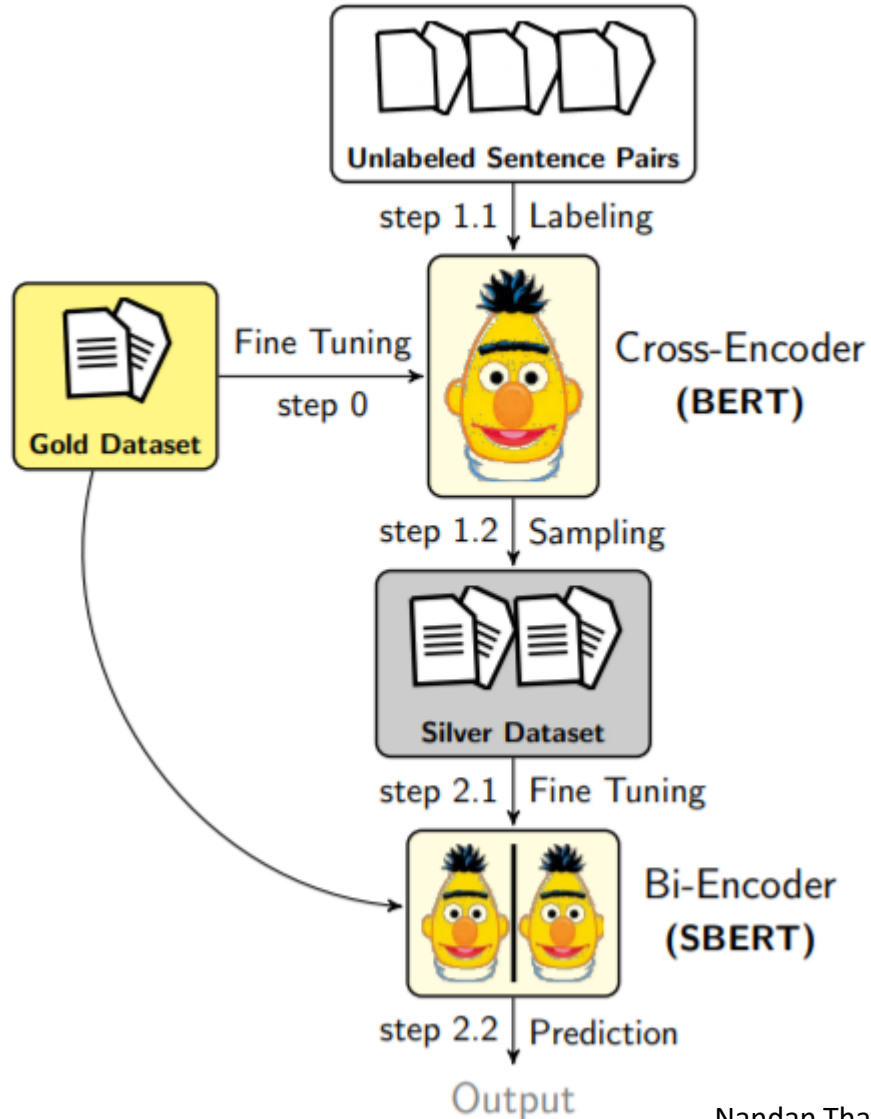
Nils Reimers

www.SBERT.net

Bi- vs Cross-Encoders – Data Efficiency

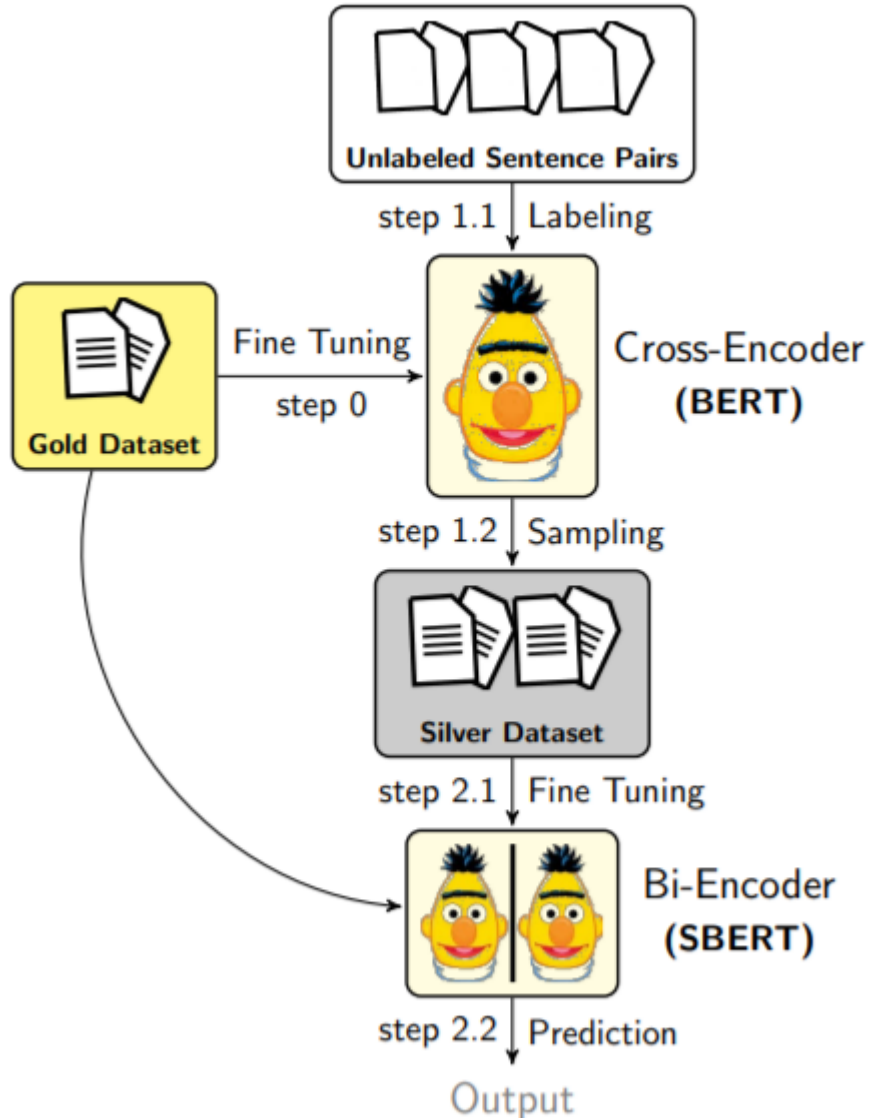


Augmented SBERT




- Train Cross-Encoder
- Sample new sentence pairs
 - Label pairs with Cross-Encoder
- Train Bi-Encoder

Augmented SBERT - Sampling



- Random sampling yields mainly pairs with low similarity:



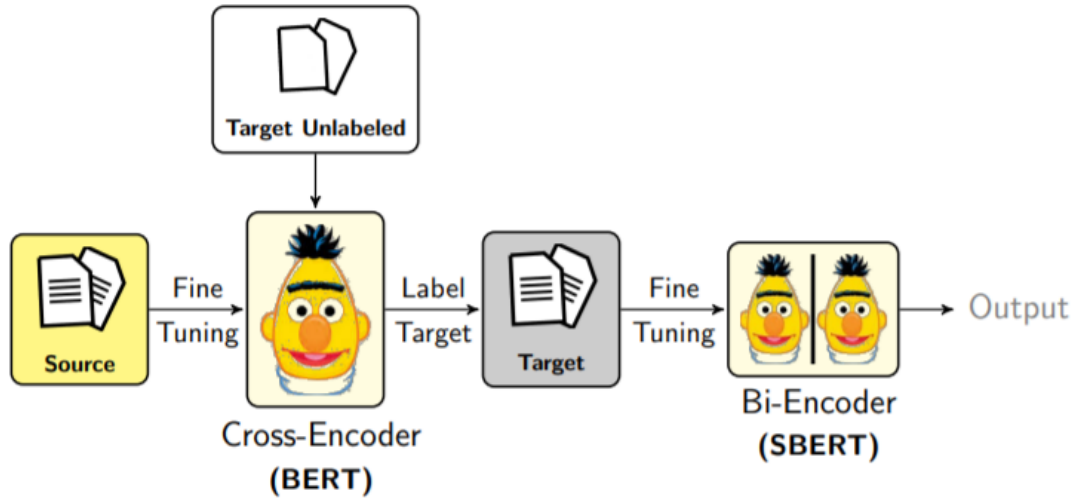
- BM25 sampling & Semantic Search yields best results

Augmented SBERT – In-Domain Results

Model	Spanish-STs	Argument Similarity	Duplicate Questions	Paraphrase Identification
BERT (Cross-Encoder)	77.5	65.1	80.4	89.0
SBERT (Bi-Encoder)	68.4	58.0	73.4	84.4
AugSBERT (Bi-Encoder)	75.1	61.5	79.3	85.5

In-domain improvement up to 6.7 points

Augmented SBERT – Cross-Domain



- Labeled training data on source domain only (Quora)
- Evaluation on specialized domains

Eval-Dataset	SBERT (Quora)	AugSBERT
AskUbuntu	50.1	60.2
Sprint	50.5	87.5
SuperUser	50.4	64.5

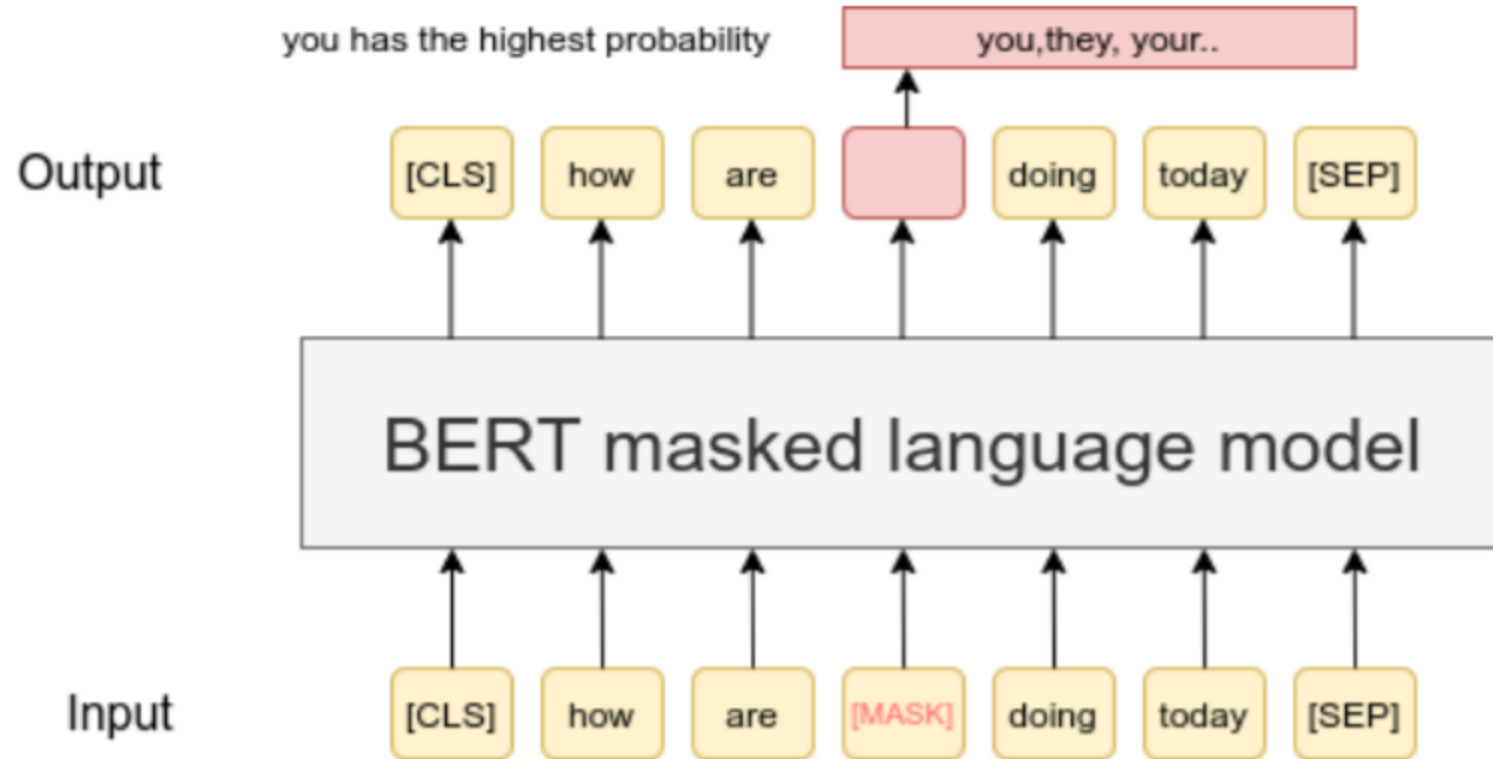
Cross-domain improvement up to 37 points

Unsupervised Embedding Methods

Nils Reimers

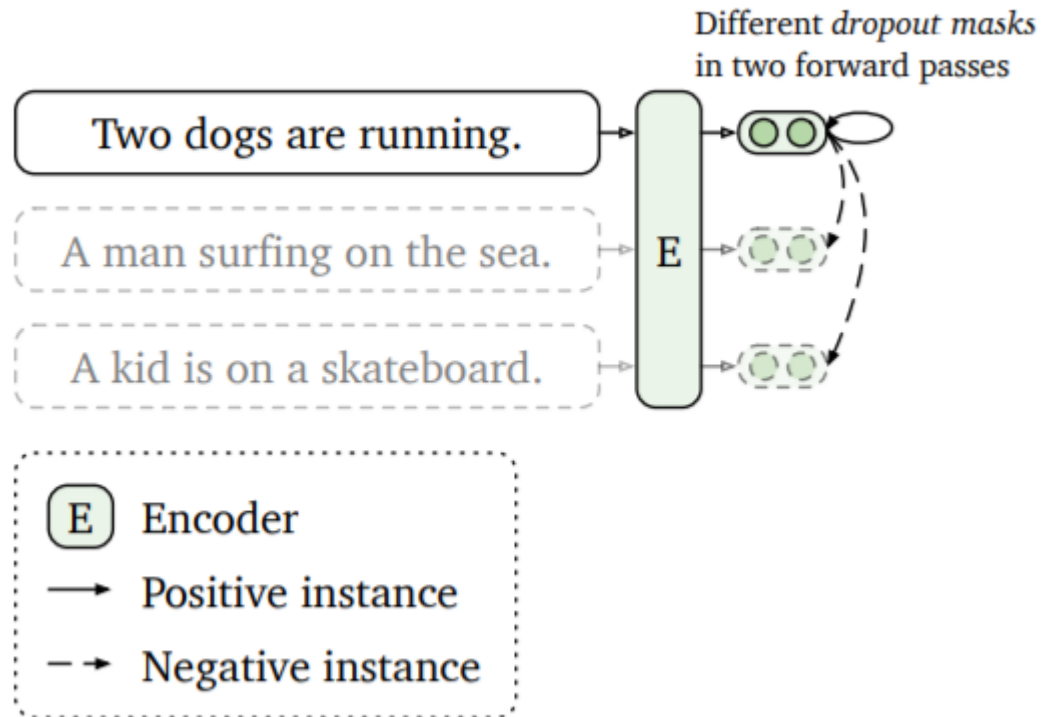
www.SBERT.net

Masked Language Model (MLM)




SimCSE

(a) Unsupervised SimCSE




- Usage of MultipleNegativeRankingLoss
- Input pairs:
 - (sent1, sent1)
 - (sent2, sent2)
 -
- Due to dropout: slightly different embeddings for $f(\text{sent1})$ and $f(\text{sent1})$

Contrastive Tension (CT)



- Initialize with two identical models
- Pass pairs with identical and with different sentences
- Maximize dot-score for identical sentences
- Minimize dot-score for different sentences

TSDAE



- Delete randomly words in the text
- Pass through the encoder
- Apply pooling to get fixed-sized text embedding
- Decoder must reconstruct text without noise from this text embedding

Issues in the Evaluation

- So far unsupervised methods evaluated on STS data
- Extremely bad way to evaluate unsupervised methods on STS datasets
 - Performance has near zero correlation to performance on real-world task
 - Simple sentences without domain specific knowledge
 - Unrealistic label distribution
- In TSDAE: Evaluation on domain specific datasets
 - AskUbuntu, StackExchange, Twitter, Scientific Publications

Evaluation


Method	Avg. over 4 datasets
TSDAE	55.2
MLM	52.9
CT	52.4
SimCSE	50.6
<i>Out-of-the-box model</i>	
SBERT on NLI+STSb	52.3

How good are unsupervised methods?

	AskUbuntu	Twitter Paraphrases	StackExchange	SciDocs
Unsupervised in-domain TSDAE on bert-base	55.6	74.1	36.2	74.5
Supervised out-of-domain mpnet + NLI + STSb	56.0	78.9	35.7	71.4
Supervised out-of-domain distilbert + MS MARCO	56.1	74.6	40.3	70.8

- Supervised pre-trained models hard to beat
- Diversity of pre-training dataset critical
 - Large, diverse dataset => great results across tasks

Unsupervised Method for Pre-Training



(d) SciDocs

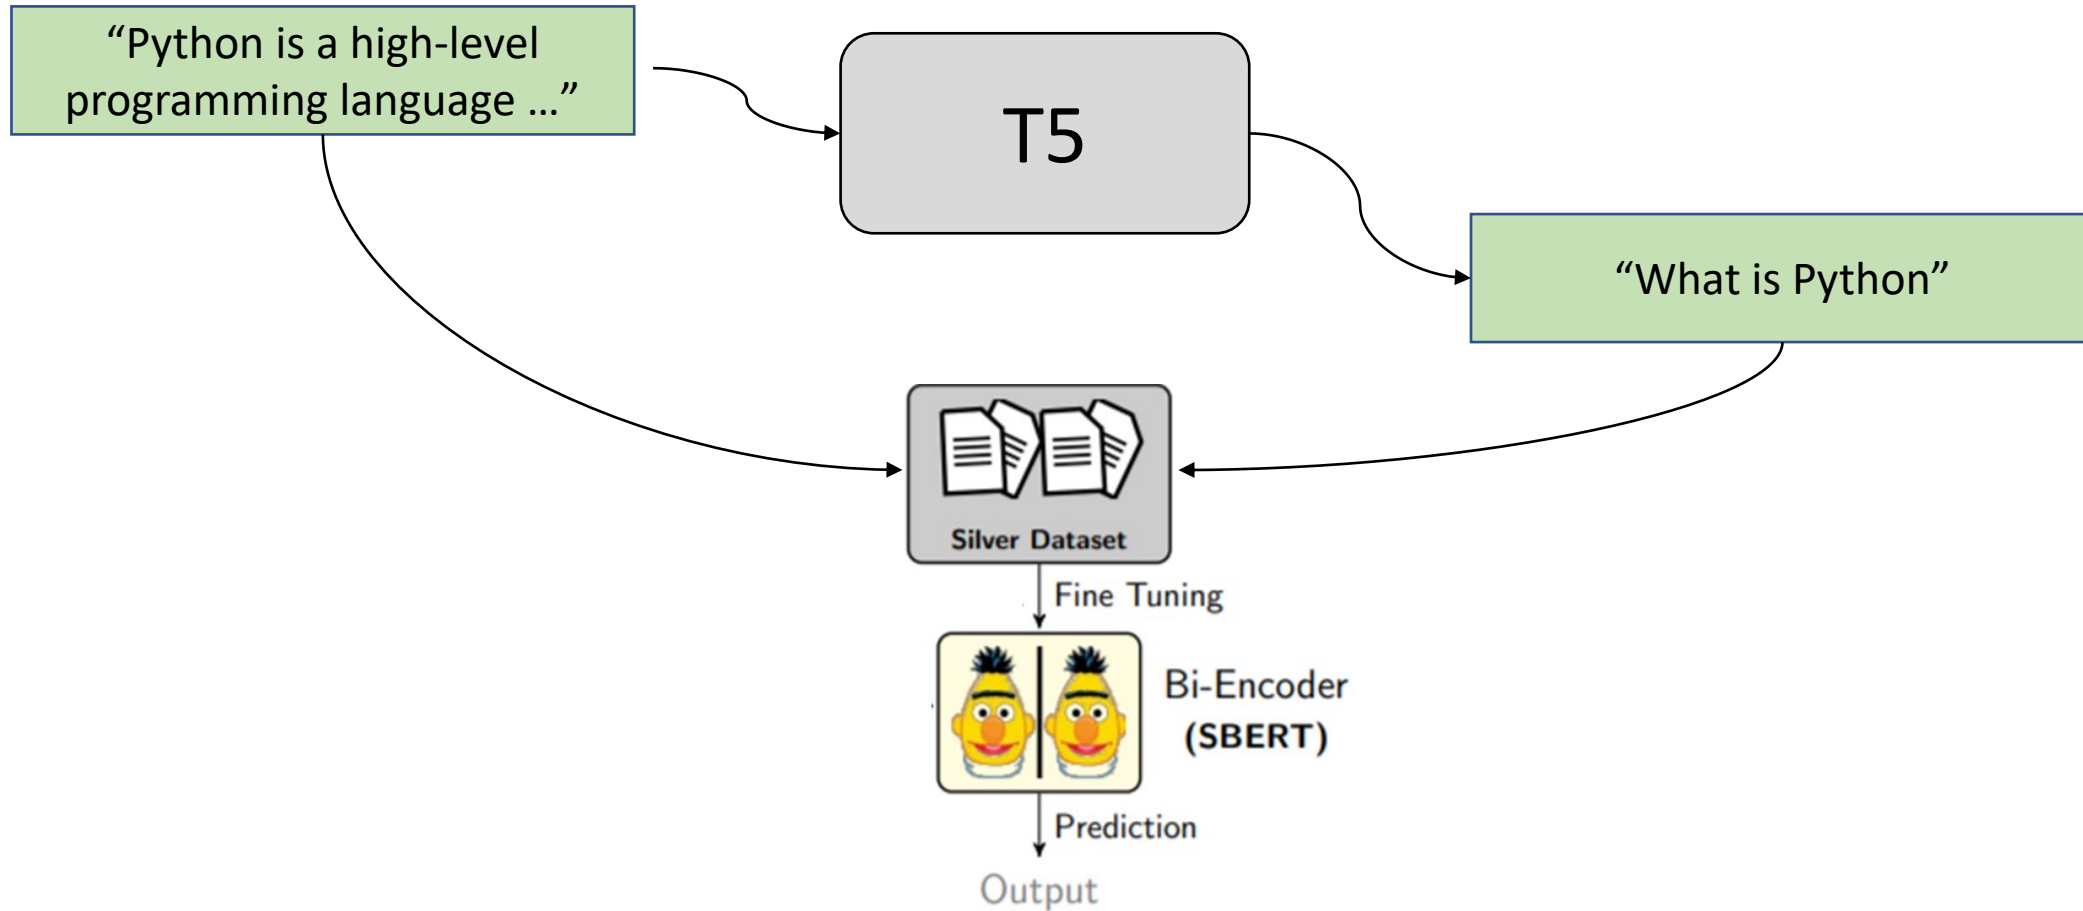
- Train unsupervised on large corpus from your domain
- Train supervised with some labels from your domain
- SimCSE / CT: Not helpful
- TSDAE / MLM: Big improvement

Domain Adaptation

Method	Unsupervised	NLI+STS -> Unsupervised	Unsupervised -> NLI+STS
TSDAE	55.2	54.2	56.5
MLM	52.9	51.1	55.9
CT	52.4	52.9	53.0
SimCSE	50.6	51.2	52.4

- First train unsupervised on your domain
- Then train supervised on available training data from other domains

GenQ: Synthetic Query Generation





Beir
Benchmarking IR

<https://github.com/UKPLab/beir>


Types of Search – BM25

How are you?
[0, 0, 0.3, 0, 0, 0, 0.1, 0, 0.2, ...]
How are you

$$\text{score}(D, Q) = \sum_{i=1}^n \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgtl}}\right)}$$


- Strong baseline
- Advantages:
 - Fast
 - Small index
 - No training required
- Disadvantages:
 - Lexical gap
 - Word order not preserved

Types of Search – Dense Retrieval




- Maps query & docs to vector spaces
- Advantages:
 - Fast
 - Overcomes lexical gap
 - Can deal with word order
- Disadvantages:
 - Require training data

Types of Search- Re-Ranking



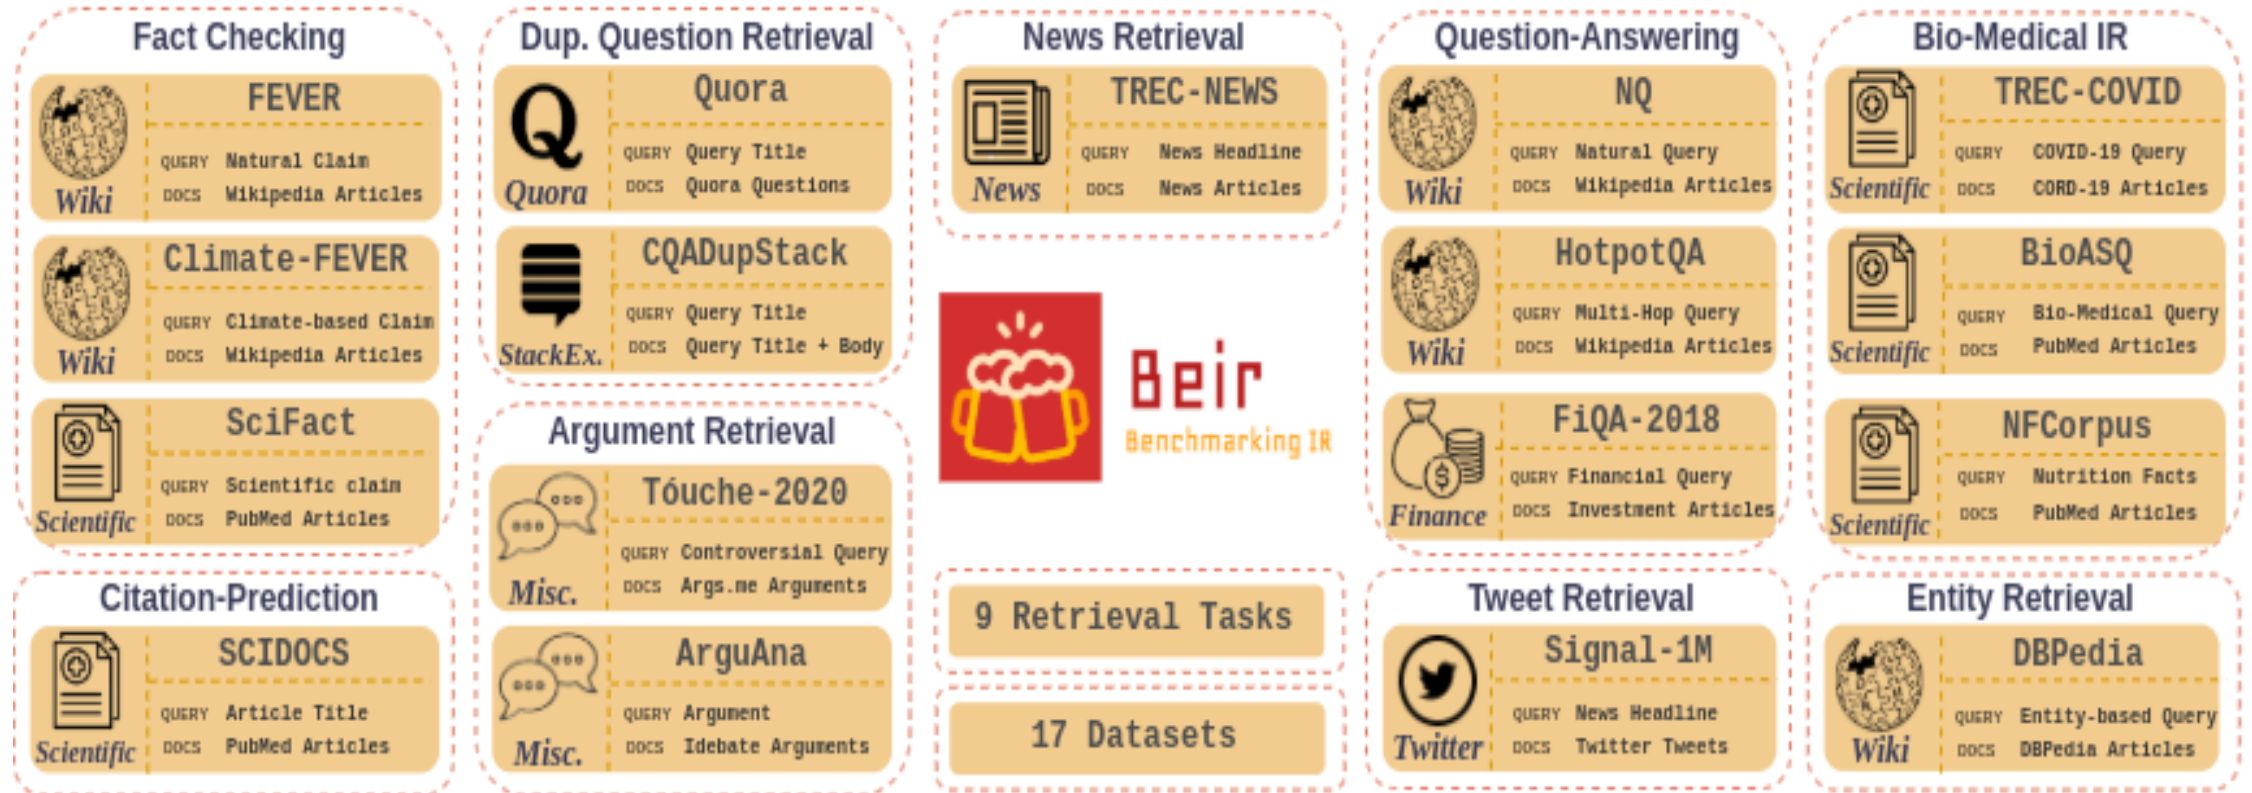
- Uses a Cross-Encoder to re-rank candidates
- Advantages:
 - High accuracy
- Disadvantages:
 - Slow
 - Requires training data

Types of Search- ColBERT




- Maps query & docs to token embeddings
- Advantages:
 - Good accuracy
- Disadvantages:
 - Large index size
 - Slow
 - Requires training data

Diverse Zero-Shot Evaluation of Neural Retrieval Methods



- How do neural retrieval models perform without in-domain training data?

Which models generalize well?



BM25 (Lexical)

BM25 is an overall strong baseline for retrieval. It doesn't require to be fine-tuned for any dataset.

ColBERT, BM25 + CE (Rerank)

Reranking Models overall generalize well across diverse domains, and outperform BM25 on retrieval tasks

ANCE, GenQ, SBERT.. (Dense)

Dense models overall are unable to generalize. They perform well with a huge source overlap with target domain

Speed – Quality – Memory-Trade-off

DBPedia (1 Million)			Retrieval Latency		Index
Rank	Model	Dim.	GPU	CPU	Size
(1)	BM25+CE	768	550ms	7100ms	0.4GB
(2)	ColBERT	128	350ms	–	20GB
(3)	BM25	–	–	20ms	0.4GB
(4)	GenQ	768	14ms	125ms	3GB
(5)	ANCE	768	20ms	275ms	3GB
(6)	SBERT	768	14ms	125ms	3GB
(7)	SPARTA	2000	–	20ms	12GB
(8)	USE-QA	512	35ms	75ms	2GB
(9)	DPR	768	19ms	230ms	3GB

- BM25+CE & ColBERT best systems
 - Slow / large index
- SBERT/ANCE better than BM25 if task similar to training data
 - Larger index (3GB vs. 0.4GB)
 - Slower on CPU (125ms vs 20ms)

Conclusion

- Multilingual Knowledge Distillation
- Cross-Encoder useful to augment Bi-Encoder with additional data
- Unsupervised embeddings
 - Challenging
 - Performance not matching supervised methods
- Information retrieval
 - Trade-off: Memory-Speed-Quality