

Deep Learning for Natural Language Processing

Lecture 10 – Text classification 4: self-attention and BERT

Dr. Martin Tutek

June 20, 2023

Ubiquitous Knowledge Processing
Department of Computer Science
Technical University of Darmstadt




[UKP Web](#)

Recap

In the previous lecture we:

- Introduced the Transformer architecture
- Explained what we gain from *contextualized* representations
- Analyzed the Transformer attention block
- Introduced byte-pair encodings & what we gain by them
- Introduced positional embeddings & why we need them

Recap: the Transformer architecture



Recap: Transformer for machine translation


We studied the Transformer encoder-decoder for machine translation



Image source: [The illustrated Transformer](#)

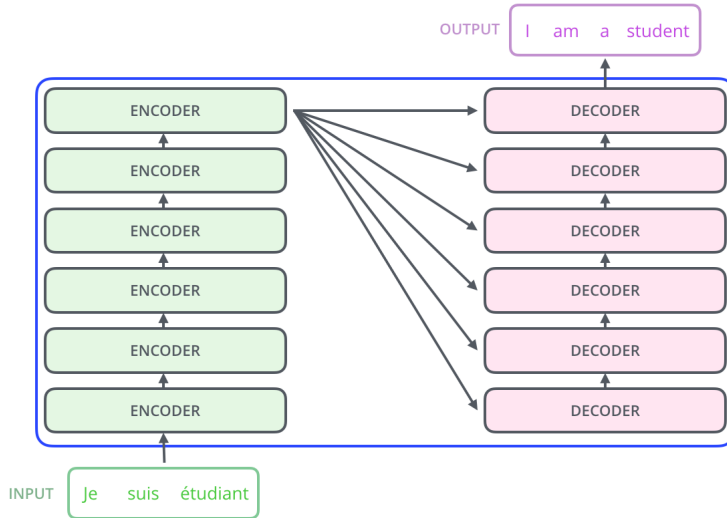
The model is built of stacked encoder and decoder blocks

Recap: Transformer encoder-decoder




The encoder and decoder are made of stacked **encoder** and **decoder blocks**

Recap: Transformer encoder-decoder

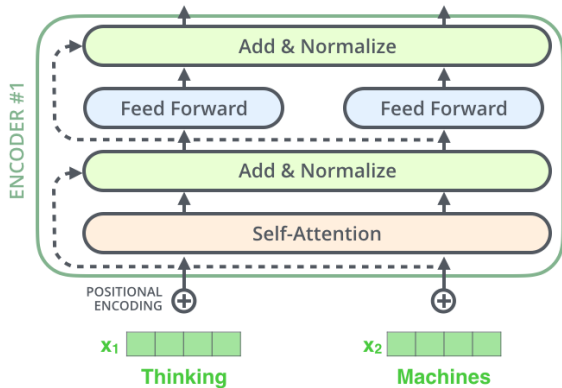


Recap: Transformer encoder and decoder block

The encoder and decoder blocks are different – the decoder block has an additional **encoder-decoder attention** (cross-attention) layer



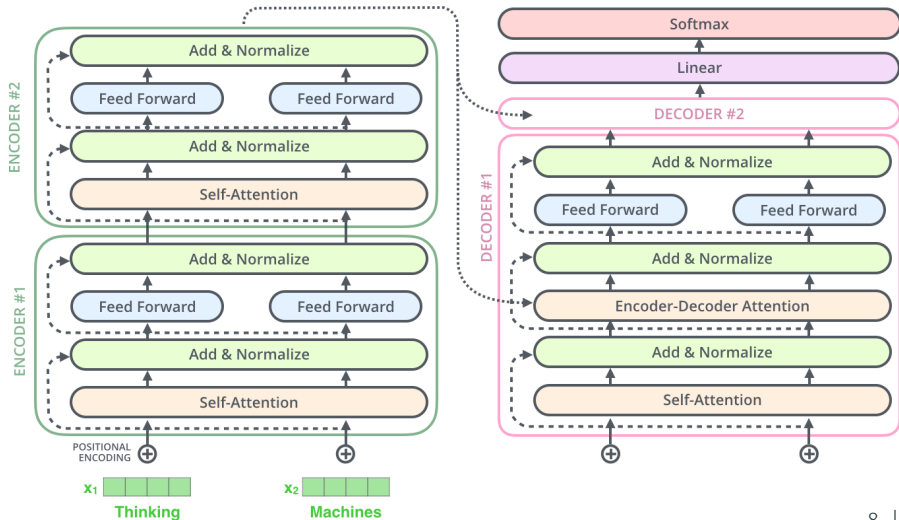
Recap: Transformer encoder block



Each Transformer encoder block consists of:

1. Self-attention: **contextualize** representations
2. **Residual** connection + **normalization**
3. **Feed-forward** layer (1 hidden layer NN)
4. **Residual** connection + **normalization**

Recap: Transformer encoder-decoder



Transformer decoding gif

Next steps?

Next steps?

BERT

Fine-tuning BERT

Variants of pretraining tasks

Motivation

The transformer encoder-decoder model is **really** good at sequence-to-sequence tasks

- It **scales** well (to many layers & parameters)
- It **performs** well on long sequences
- It's **easier to optimize** (residual connections)
- It's **faster to run** (parallel processing in encoder)

However – we can **only** use it for the **task it was trained on**

Motivation

Transfer learning

... is applying knowledge gained when **solving one task** to a **related** task.

Gained knowledge → encoded in a **trained model**.

Where have we seen something like this before?

- **Word2vec** (CBOW & Skip-gram)

We **train** the word embeddings on an *auxiliary task*, then use them as input for other models Can we apply this to **Transformer** encoders to obtain pretrained **contextualized** embeddings?

BERT

Next steps?

BERT

- Fine-tuning BERT

- Variants of pretraining tasks



BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova
Google AI Language

What **we have**

- A model: the Transformer

What **we want**

- Pretrained *contextualized* word representations


What **we need**

- The **auxiliary pretraining task**
- ... ideally, it should not require labeled data (expensive)

BERT: pretraining objective

Recall: what was the word2vec training objective?

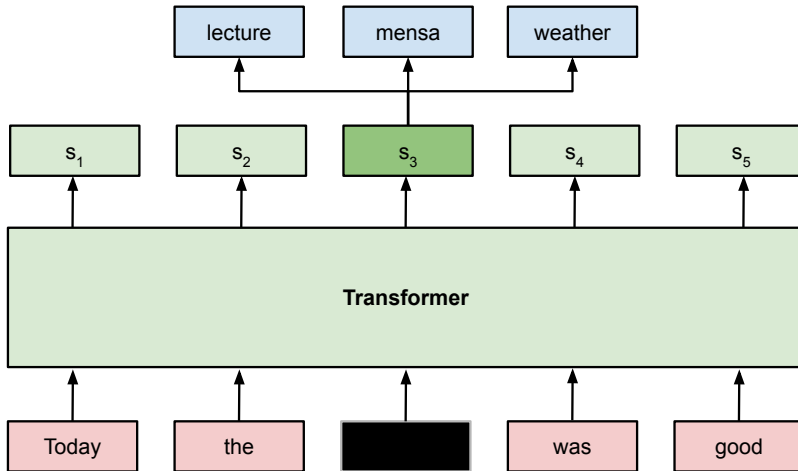
- **CBOW:** predict **center word** given **context words**



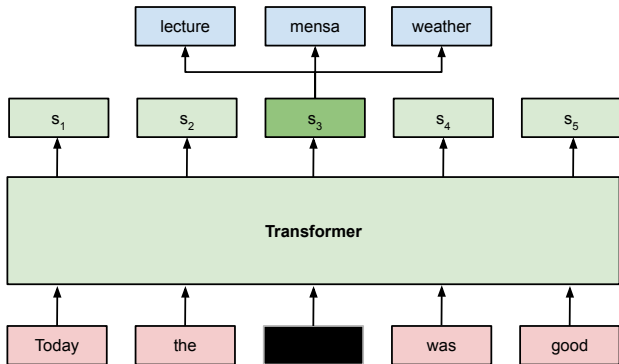
- **Skip-gram:** predict **context words** given **center word**

Can we use a **similar** task with Transformer models?

BERT: pretraining objective



BERT: pretraining objective



1. For an input sequence $\{x_i\}_{i=1}^n$, we **mask** an input token(s)
2. We encode the inputs with a Transformer **encoder**
3. We **reconstruct** the masked token(s)

What is *masking*?

BERT: masked language modeling (MLM)

The **Transformer** model contextualizes an input sequence $\{\mathbf{x}_i\}_{i=1}^n$ of (subword) tokens into a sequence of hidden states $\{\mathbf{s}_i\}_{i=1}^n$.

1. With probability p_{mlm} , mask **each input token** ($p_{mlm} = 0.15$)
2. **If** a token is masked
 - 80% of the time, replace it with a special **<MASK>** token
 - 10% of the time, replace it with a **random** token
 - 10% of the time, **do not mask it**
3. **Only** for the **masked tokens**
 - Predict which token was masked


BERT: masked language modeling (MLM) [Image source]

Use the output of the masked word's position to predict the masked word

Possible classes:
All English words

0.1%	Aardvark
...	...
10%	Improvisation
...	...
0%	Zyzyva

FFNN + Softmax



Randomly mask
15% of tokens

Input

[CLS] Let's stick to improvisation in this skit

BERT: masked language modeling (MLM)

And... it works

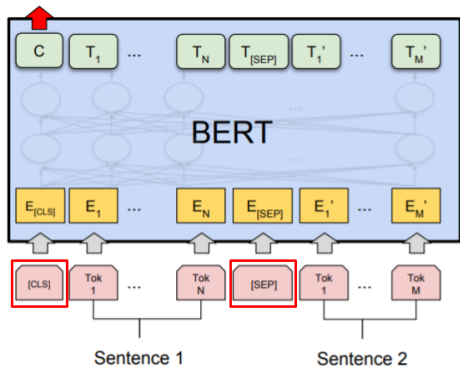
System	MNLI-(m/mm) 392k	QQP 363k	QNLI 108k	SST-2 67k	CoLA 8.5k	STS-B 5.7k	MRPC 3.5k	RTE 2.5k	Average -
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

GLUE Test results (GLUE benchmark). Table from BERT paper

However – we are *not there yet*

- How do we obtain **sentence representations**?

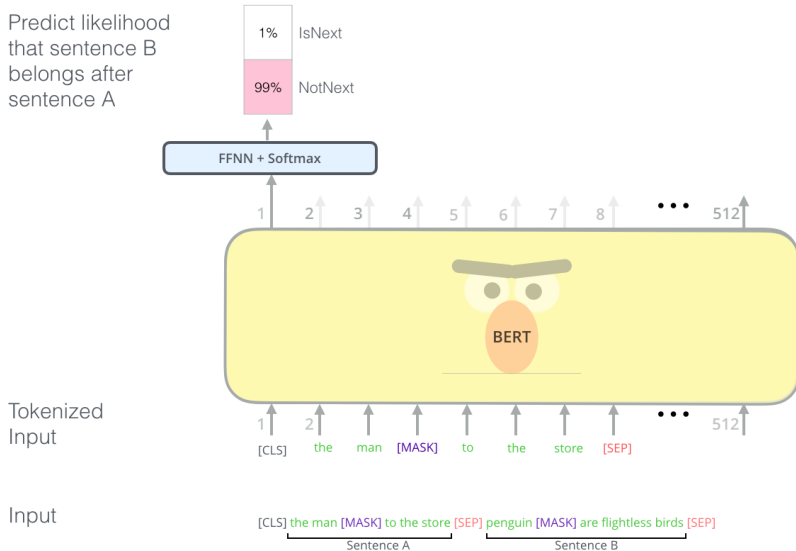
BERT: next sentence prediction (NSP)



We want to have a sentence representation **out-of-the-box**

1. We add a special CLS token as the **sentence representation**
2. We add a special SEP token to separate **two input sentences**
3. We predict (based on the CLS token) if Sent.1 **directly precedes** Sent.2 in the pretraining dataset

BERT: next sentence prediction (NSP) [Image source]



BERT: pretraining

Combining the **MLM** and **NSP** objectives:

1. Take a **large corpus** of unstructured text (Wikipedia, BookCorpus) and retain information about position of **sentences** in each article
2. Create the input sequence: $\mathbf{in} = [\text{CLS}]\{\mathbf{x}_i^1\}_{i=1}^{n_1}[\text{SEP}]\{\mathbf{x}_i^2\}_{i=1}^{n_2}$
 - 2.1 Sample Sentence_1 from the dataset
 - 2.2 With $p_{nsp} = 0.5$, take Sentence_2 as the following sentence (heads) **or** randomly sample it (tails)
3. Mask $p_{mlm} = 0.15$ of **non-special** input tokens (recall: 80/10/10)
4. Encode inputs with transformer encoder: $\mathbf{trf}(\mathbf{in}) \rightarrow \{\mathbf{s}_i\}_{i=1}^{n_1+n_2+2}$
5. Pretraining tasks
 - 5.1 **MLM**: reconstruct masked tokens $\mathbf{s}_i \rightarrow \mathbf{x}_i \quad \forall i \in \{\text{masked}\}$;
 - 5.2 **NSP**: predict if sentences are successors $\mathbf{s}_1 \rightarrow \{0, 1\}$.

BERT is a **pretrained language model (PLM)**

- Through a language modeling pretraining task, the model has learned to recognize **patterns of language** and apply them for the task of **text reconstruction**
- Text reconstruction is, however, **rarely useful** in isolation
- In practice: use the PLM as a **starting point** (a very good initialization) for **fine-tuning** (additional training) for another task

BERT: applications

Now we have the pretrained BERT – a Transformer encoder.



What's next?

- How to use this model for **downstream tasks**?

BERT

Fine-tuning BERT

Using BERT for NLP tasks

Fine-tuning is the procedure where we start from a base **pretrained** model and adapt its internal representations to our task.

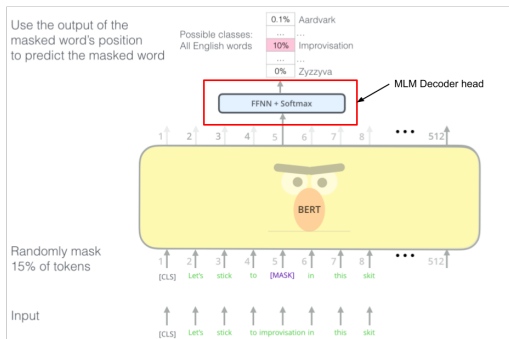
Fine-tuning variants:

1. **Vanilla fine-tuning:** add a **decoder head** to the model, then:
 - 1.1 Train **only** the decoder head;
 - 1.2 Train **progressively** more layers: first the decoder head, then also the last layer, then also the second to last ...;
 - 1.3 Train the **entire network** at the same time, maybe with **different learning rates** per layer.
2. **Adapters:** additional randomly initialized layers inserted inside the transformer layers
3. **Prompting & in-context learning:** future lectures

Using BERT for NLP tasks

What exactly are **decoder heads**?

Randomly initialized additional layers (usually linear) added **on top** of the pretrained model which **perform the downstream task**.



Using BERT: single sequence classification

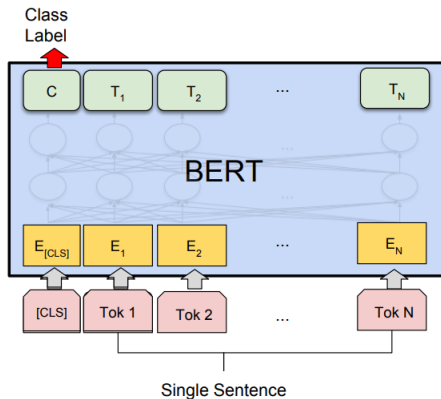


Image from BERT paper

For single **sequence** classification:

1. Add a randomly initialized **task decoder head** to the model
2. **Encode** the sequence along with the [CLS] special token
3. Use the **[CLS]** representation as input to decoder head

Alternatives to using CLS?

- Averaging over **token representations**
- ... from the **last four layers**

Using BERT: sentence pair classification

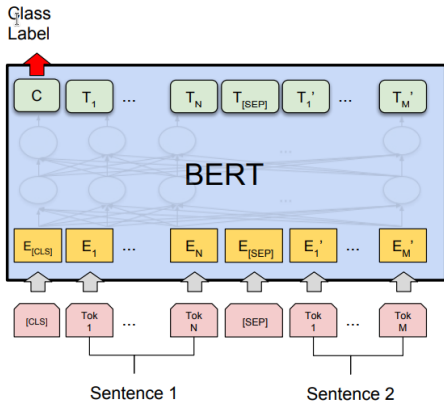


Image from BERT paper

For **pair sequence** classification:

1. Add a randomly initialized **task decoder head** to the model
2. **Encode both sequences** along with the [CLS] special token
3. Use the **[CLS]** token representation as input to decoder head

Using BERT: span extraction QA

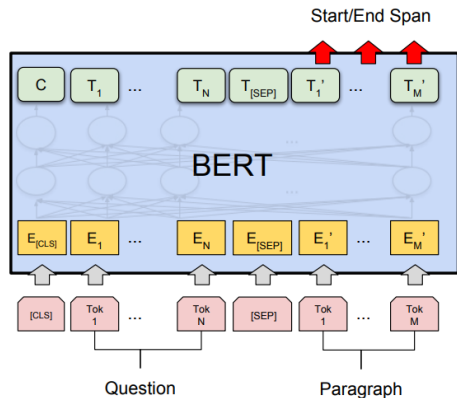


Image from BERT paper

For **span extraction QA**:

1. Add randomly initialized **start-of-span** and **end-of-span vectors** to the model.
2. **Encode both sequences**
3. Highest dot product of **token representation** with start-of-span and end-of-span is the predicted span
 - ... such that end-of-span > start-of-span

Using BERT: sequence labeling

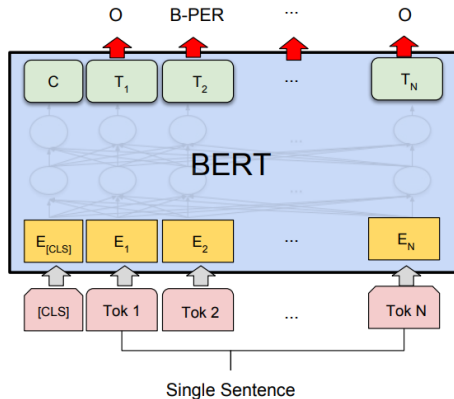


Image from BERT paper

For **sequence labeling**:

1. Add a randomly initialized **task decoder head** to the model
2. **Encode** the sequence along with the [CLS] special token
3. Use the **token representations** as inputs to decoder head

BERT

Variants of pretraining tasks

Variants of pretraining tasks

We have used **MLM** and **NSP** – are there some *better* tasks for pretraining language models?

Perhaps in an **encoder-decoder** setup?




Image from BART paper

Variants of pretraining tasks

- Language modeling
- Token masking: MLM
- **Token deletion**: masking, but **completely removes tokens** from input – model needs to determine where a token is missing
- **Text infilling**: masking, but **multiple** tokens are replaced with a **single** [MASK] token at the same time
- **Sentence permutation**: input **permuted sentence**, reconstruct correct word order (*linearization*)
- **Document rotation**: document is rotated so that it starts from a **random token**. The model has to determine the actual start of the document.

Variants of (supervised) pretraining tasks

What if we decide to use **supervised data** – but from various datasets?





Image from T5 paper

Pretrained language model architectures



(a) BERT: Random tokens are replaced with masks, and the document is encoded bidirectionally. Missing tokens are predicted independently, so BERT cannot easily be used for generation.

(b) GPT: Tokens are predicted auto-regressively, meaning GPT can be used for generation. However words can only condition on leftward context, so it cannot learn bidirectional interactions.



(c) BART: Inputs to the encoder need not be aligned with decoder outputs, allowing arbitrary noise transformations. Here, a document has been corrupted by replacing spans of text with mask symbols. The corrupted document (left) is encoded with a bidirectional model, and then the likelihood of the original document (right) is calculated with an autoregressive decoder. For fine-tuning, an uncorrupted document is input to both the encoder and decoder, and we use representations from the final hidden state of the decoder.

Takeaways

- BERT is a pretrained language model which produces **contextualized** token representations of input text
- It can be used as an **initialization** (starting point) for **fine-tuning** task-specific models
 - Extras: [CLS] and [SEP] tokens
 - Applications of BERT in classification, sequence labeling and span-extraction QA
- Other pretraining tasks are also viable
 - **Unsupervised**: sentence permutation, text infilling
 - **Supervised**: translation, summarization

Useful resources

- [The annotated Transformer](#) by Sasha Rush
- [The illustrated Transformer](#) by Jay Allamar
- [The illustrated BERT](#) by Jay Allamar

License and credits

Licensed under Creative Commons
Attribution-ShareAlike 4.0 International (CC BY-SA 4.0)



Credits

Martin Tutek

Content from ACL Anthology papers licensed under CC-BY <https://www.aclweb.org/anthology>